

Type *Markdown* and LaTeX: α^2

Ninjacart : Image classification using CNN and transfer learning

In []:

Link - <https://colab.research.google.com/drive/1A6jRmhFRFHIFiUYou7HDzlwT6F6MC8bk?usp=sharing>

Summary **and** Insights :

1. CNN **from** Scratch: 14,21,148 params

Validation accuracy - 93.38%

Test Accuracy - 88.03

2. VGG : 3,16,71,740 params

Validation accuracy - 96.12%

Test Accuracy - 87.18 %

3. Resnet : 7,51,21,020

Validation accuracy - 98.00%%

Test Accuracy - 90.88%

4. Mobilenet : 1,27,55,900

Validation accuracy - 98.25%

Test Accuracy - 92.31%

The best model **is** Mobilenet **as** it has the least number of parameters **and** the best testing accuracy.

To train a CNN we will need a lot of images. Tensorboard showed **in** the last page of the notebook.

In []:

```
!gdown 1clZX-lV_MLxKHSyeyTheX50CQtNCUcqT
# Downloading the Data
```

Downloading...

From: https://drive.google.com/uc?id=1clZX-lV_MLxKHSyeyTheX50CQtNCUcqT (https://drive.google.com/uc?id=1clZX-lV_MLxKHSyeyTheX50CQtNCUcqT)

To: /content/ninjacart_data.zip

100% 275M/275M [00:08<00:00, 31.4MB/s]

In []:

```
import shutil
shutil.rmtree('ninjacart_data')
```


In []:

```

# Import common libraries
import os
import glob
import random
import numpy as np
import pandas as pd
import sklearn.metrics as metrics
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import tensorflow as tf
from tensorflow.keras.preprocessing.image import load_img, img_to_array, random_shift
from tensorflow.keras.preprocessing.image import array_to_img

import tensorflow as tf
from tensorflow import keras # this allows <keras.> instead of <tf.keras.>
from tensorflow.keras import layers # this allows <Layers.> instead of <tf.keras.Layers.>
tf.keras.utils.set_random_seed(111) # set random seed

import os
import numpy as np
import tensorflow as tf
tf.keras.utils.set_random_seed(111)
from tensorflow import keras
from tensorflow.keras import layers, regularizers
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.metrics as metrics

import keras
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Flatten, Conv2D, MaxPooling2D
from keras.layers import BatchNormalization
import numpy as np
from tensorflow.keras.callbacks import TensorBoard

# To suppress any warnings during the flow
import warnings
warnings.filterwarnings('ignore')

```

In []:

```

class_dirs = os.listdir("ninjacart_data/train") # List all directories inside "train" folder
image_dict = {} # dict to store image array(key) for every class(value)
count_dict = {} # dict to store count of files(key) for every class(value)
# iterate over all class_dirs
for cls in class_dirs:
    # get list of all paths inside the subdirectory
    file_paths = glob.glob(f'ninjacart_data/train/{cls}/*')
    # count number of files in each class and add it to count_dict
    count_dict[cls] = len(file_paths)
    # select random item from list of image paths
    image_path = random.choice(file_paths)
    # load image using keras utility function and save it in image_dict
    image_dict[cls] = tf.keras.utils.load_img(image_path)

```

In []:

```
plt.figure(figsize=(15, 12))  
# iterate over dictionary items (class label, image array)  
for i, (cls,img) in enumerate(image_dict.items()):  
    # create a subplot axis  
    ax = plt.subplot(3, 4, i + 1)  
    # plot each image  
    plt.imshow(img)  
    # set "class name" along with "image size" as title  
    plt.title(f'{cls}, {img.size}')  
    plt.axis("off")  
# Plotting a few images from each class for viewing  
# We can conclude that image sizes are not constant
```

onion, (244, 207)



tomato, (400, 500)



indian market, (870, 580)



potato, (267, 188)



In []:

```
## Let's now Plot the Data Distribution of Training Data across Classes
df_count_train = pd.DataFrame({
    "class": count_dict.keys(),    # keys of count_dict are class labels
    "count": count_dict.values(),  # value of count_dict contain counts of each class
})
print("Count of training samples per class:\n", df_count_train)

# draw a bar plot using pandas in-built plotting function
df_count_train.plot.bar(x='class', y='count', title="Training Data Count per class")

# Below table shows the number of samples available in each class
# Data is imbalanced
# There are a total of 4 classes
```

Count of training samples per class:

	class	count
0	onion	849
1	tomato	789
2	indian market	599
3	potato	898

Out[7]:

<Axes: title={'center': 'Training Data Count per class'}, xlabel='class'>



Data augmentation to make sample size as 1000 for each class

1. Random Crop
2. Image rotation
3. Random flip

In []:

```
class_dirs = os.listdir("ninjacart_data/train") # list all directories inside "train" folder
image_dict = {} # dict to store image array(key) for every class(value)
count_dict = {} # dict to store count of files(key) for every class(value)
# iterate over all class_dirs
for cls in class_dirs:
    # get list of all paths inside the subdirectory
    file_paths = glob.glob(f'ninjacart_data/train/{cls}/*')
    last=len(file_paths)
    list_path=random.choices(file_paths,k=1000-last)
    print(len(list_path))
    for i in range(0,len(list_path)-2,3):
        #print(i,cls)

        img = load_img(list_path[i],target_size=(300,300))
        x = img_to_array(img)
        x_crop = tf.image.random_crop(value=x, size=(224, 224, 3))
        array_to_img(x_crop)
        out_filename = f'{str(i)}.jpg'
        img.save(f'ninjacart_data/train/{cls}/' + out_filename)

        img = load_img(list_path[i+1],target_size=(300,300))
        x = img_to_array(img)
        x_rot=tf.keras.preprocessing.image.random_rotation(x, rg=125,channel_axis = 2, row_axis=0)
        array_to_img(x_rot)
        out_filename = f'{str(i+1)}.jpg'
        img.save(f'ninjacart_data/train/{cls}/'+ out_filename)

        img = load_img(list_path[i+2],target_size=(300,300))
        x = img_to_array(img)
        x_flip = tf.image.random_flip_left_right(x, 10)
        array_to_img(x_flip)
        out_filename = f'{str(i+2)}.jpg'
        img.save(f'ninjacart_data/train/{cls}/' + out_filename)
```

151
211
401
102

In []:

```

class_dirs = os.listdir("ninjacart_data/train") # List all directories inside "train" folder
image_dict = {} # dict to store image array(key) for every class(value)
count_dict = {} # dict to store count of files(key) for every class(value)
# iterate over all class_dirs
for cls in class_dirs:
    # get list of all paths inside the subdirectory
    file_paths = glob.glob(f'ninjacart_data/train/{cls}/*')
    # count number of files in each class and add it to count_dict
    count_dict[cls] = len(file_paths)
    # select random item from list of image paths
    image_path = random.choice(file_paths)
    # load image using keras utility function and save it in image_dict
    image_dict[cls] = tf.keras.utils.load_img(image_path)

```

In []:

```

plt.figure(figsize=(15, 12))
# iterate over dictionary items (class label, image array)
for i, (cls,img) in enumerate(image_dict.items()):
    # create a subplot axis
    ax = plt.subplot(3, 4, i + 1)
    # plot each image
    plt.imshow(img)
    # set "class name" along with "image size" as title
    plt.title(f'{cls}, {img.size}')
    plt.axis("off")

# Checking image in each class

```



In []:

```
## Let's now Plot the Data Distribution of Training Data across Classes
df_count_train = pd.DataFrame({
    "class": count_dict.keys(),      # keys of count_dict are class labels
    "count": count_dict.values(),    # value of count_dict contain counts of each class
})
print("Count of training samples per class:\n", df_count_train)

# draw a bar plot using pandas in-built plotting function
df_count_train.plot.bar(x='class', y='count', title="Training Data Count per class")

# All classes have image count close to 1000
# Now it is a balanced dataset
```

Count of training samples per class:

	class	count
0	onion	999
1	tomato	999
2	indian market	998
3	potato	1000

Out[11]:

<Axes: title={'center': 'Training Data Count per class'}, xlabel='class'>



Creating Validation Data Set of size 200 each

In []:

```

import os
import random
from shutil import move

train_dir = 'ninjaart_data/train'
validation_dir = 'ninjaart_data/validation'

# Create the validation directory if it does not exist
if not os.path.exists(validation_dir):
    os.mkdir(validation_dir)

# Get a List of all the files in the train directory
files = os.listdir(train_dir)
for i in files:
    os.mkdir(f'ninjaart_data/validation/{i}')

for i in files:
    files = os.listdir(f'ninjaart_data/train/{i}')
    #print(len(files))
    random_files = random.sample(files, 200)
    #print(len(files))

# Move the selected files to the validation directory
for file in random_files:
    move(f'ninjaart_data/train/{i}/{file}', f'ninjaart_data/validation/{i}/{file}')

```

In []:

```

class_dirs = os.listdir("ninjaart_data/train") # List all directories inside "train" folder
image_dict = {} # dict to store image array(key) for every class(value)
count_dict = {} # dict to store count of files(key) for every class(value)
# iterate over all class_dirs
for cls in class_dirs:
    # get list of all paths inside the subdirectory
    file_paths = glob.glob(f'ninjaart_data/train/{cls}/*')
    # count number of files in each class and add it to count_dict
    count_dict[cls] = len(file_paths)

```

In []:

count_dict

Out[14]:

```
{'onion': 799, 'tomato': 799, 'indian market': 798, 'potato': 800}
```

In []:

```

class_dirs = os.listdir("ninjaart_data/validation") # List all directories inside "train" folder
image_dict = {} # dict to store image array(key) for every class(value)
count_dict = {} # dict to store count of files(key) for every class(value)
# iterate over all class_dirs
for cls in class_dirs:
    # get list of all paths inside the subdirectory
    file_paths = glob.glob(f'ninjaart_data/validation/{cls}/*')
    # count number of files in each class and add it to count_dict
    count_dict[cls] = len(file_paths)

```

In []:

```
count_dict  
# We have 200 images in validation dataset (20% from training data)
```

Out[16]:

```
{'onion': 200, 'tomato': 200, 'indian market': 200, 'potato': 200}
```

Loading Data

In []:

```
def load_data(base_dir="ninjacart_data"):  
  
    print('\nLoading Data...')  
    train_data = tf.keras.utils.image_dataset_from_directory(  
        f"{base_dir}/train", shuffle=True, label_mode='categorical'  
    )  
    val_data = tf.keras.utils.image_dataset_from_directory(  
        f"{base_dir}/validation", shuffle=False, label_mode='categorical'  
    )  
    test_data = tf.keras.utils.image_dataset_from_directory(  
        f"{base_dir}/test", shuffle=False, label_mode='categorical'  
    )  
    return train_data, val_data, test_data, train_data.class_names
```

In []:

```
train_data, val_data, test_data, class_names = load_data()
```

```
Loading Data...  
Found 3196 files belonging to 4 classes.  
Found 800 files belonging to 4 classes.  
Found 351 files belonging to 4 classes.
```

Setting image dimension and rescaling images

In []:

```
def preprocess(train_data, val_data, test_data, target_height=256, target_width=256):

    # Data Processing Stage with resizing and rescaling operations
    data_preprocess = keras.Sequential(
        name="data_preprocess",
        layers=[
            layers.Resizing(target_height, target_width),
            layers.Rescaling(1.0/255),
        ]
    )

    # Perform Data Processing on the train, val, test dataset
    train_ds = train_data.map(lambda x, y: (data_preprocess(x), y), num_parallel_calls=tf.data.AUTOTUNE)
    val_ds = val_data.map(lambda x, y: (data_preprocess(x), y), num_parallel_calls=tf.data.AUTOTUNE)
    test_ds = test_data.map(lambda x, y: (data_preprocess(x), y), num_parallel_calls=tf.data.AUTOTUNE)

    return train_ds, val_ds, test_ds
```

In []:

In []:

```
train_ds, val_ds, test_ds = preprocess(train_data, val_data, test_data)
```

In []:

```
L2Reg = tf.keras.regularizers.L2(l2=1e-6)
# Regularizers
```

In []:

```
def scheduler(epoch, lr):
    if epoch < 20:
        return lr
    if epoch < 35:
        return 0.0001
    if epoch < 50:
        return 0.00001
    if epoch < 80:
        return 0.000001
    else:
        return 0.0000001

# Learning rate scheduler
```

Function for fitting

1. Saving weights

2. Learning rate scheduler

In []:

```
def compile_train_v1(model, train_ds, val_ds, tb_callback, epochs=100, ckpt_path="/tmp/checkpoint"):
    model.compile(optimizer=tf.keras.optimizers.Adam(beta_1=0.9, beta_2=0.999),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    model_fit = model.fit(train_ds, validation_data=val_ds, epochs=epochs, batch_size=230, callbacks=[
        keras.callbacks.ModelCheckpoint(ckpt_path, save_weights_only=True, monitor='val_accuracy'),
        tf.keras.callbacks.LearningRateScheduler(scheduler), tb_callback, tf.keras.callbacks.EarlyStopping],
                          return_model_fit=True)
    return model_fit
```

Function to plot accuracy and confusion matrix

In []:

```
def print_accuracy_stats(model, ds, class_names, ckpt_path):
    model.load_weights(ckpt_path)
    true_onehot = tf.concat([y for x, y in ds], axis=0)
    true_categories = tf.argmax(true_onehot, axis=1)
    y_pred = model.predict(ds)
    predicted_categories = tf.argmax(y_pred, axis=1)

    test_acc = metrics.accuracy_score(true_categories, predicted_categories) * 100
    print(f'\n Accuracy: {test_acc:.2f}%\n')

# Note: This doesn't work with shuffled datasets
def plot_confusion_matrix(model, ds, class_names, ckpt_path):
    model.load_weights(ckpt_path)
    true_onehot = tf.concat([y for x, y in ds], axis=0)
    true_categories = tf.argmax(true_onehot, axis=1)
    y_pred = model.predict(ds)
    predicted_categories = tf.argmax(y_pred, axis=1)
    cm = metrics.confusion_matrix(true_categories, predicted_categories) # Last batch
    sns.heatmap(cm, annot=True, xticklabels=class_names, yticklabels=class_names, cmap="YlGnBu")
    plt.show()
```

In []:

CNN Classifier model from scratch

Tensorboard

In []:

```
%load_ext tensorboard
log_folder='logs/1'
%reload_ext tensorboard
!rm -rf logs
tb_callback=TensorBoard(log_dir=log_folder,histogram_freq=1)
```


In []:

```
model1 = keras.Sequential(  
    name="model_cnn",  
    layers=[  
        layers.Conv2D(filters=16, kernel_size=(3,3), padding="same", input_shape=(256,256,3),  
                        kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.MaxPooling2D(),  
  
        layers.Conv2D(filters=32, kernel_size=(5,5), padding="same",  
                        kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.MaxPooling2D(),  
  
        layers.Conv2D(filters=64, kernel_size=(6,6), padding="same",  
                        kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.MaxPooling2D(),  
  
        layers.Conv2D(filters=128, kernel_size=(7,7), padding="same",  
                        kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.MaxPooling2D(),  
  
        layers.Conv2D(filters=256, kernel_size=(8,8), padding="same",  
                        kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.MaxPooling2D(),  
  
        layers.Conv2D(filters=512, kernel_size=(8,8), padding="same",  
                        kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.MaxPooling2D(),  
  
        layers.GlobalAveragePooling2D(),  
        layers.Flatten(),  
  
        layers.Dense(units=512, kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.Dropout(0.2),  
  
        layers.Dense(units=256, kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.Dropout(0.2),  
  
        layers.Dense(units=128, kernel_regularizer = L2Reg),  
        layers.Activation("relu"),  
        layers.BatchNormalization(),  
        layers.Dropout(0.3),  
    ]  
)
```

```
layers.Dense(units=64, kernel_regularizer = L2Reg),
layers.Activation("relu"),
layers.BatchNormalization(),
layers.Dropout(0.3),

layers.Dense(units=32, kernel_regularizer = L2Reg),
layers.Activation("relu"),
layers.BatchNormalization(),
layers.Dropout(0.2),

layers.Dense(units=16, kernel_regularizer = L2Reg),
layers.Activation("relu"),
layers.BatchNormalization(),
layers.Dropout(0.2),

layers.Dense(units=8, kernel_regularizer = L2Reg),
layers.Activation("relu"),
layers.BatchNormalization(),
layers.Dropout(0.1),

layers.Dense(units=4, activation='softmax')
]
```

In []:

```
model1.summary()
```

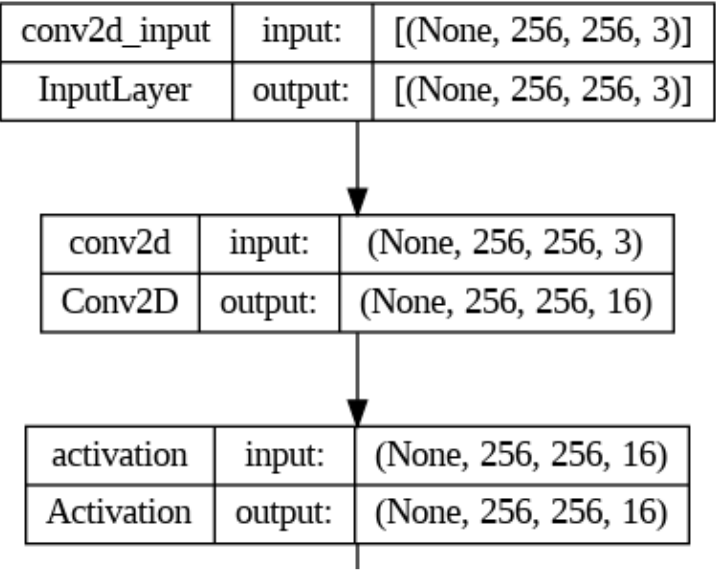
Model: "model_cnn"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 256, 256, 16)	448
activation (Activation)	(None, 256, 256, 16)	0
batch_normalization (Batch Normalization)	(None, 256, 256, 16)	64
max_pooling2d (MaxPooling2D)	(None, 128, 128, 16)	0
conv2d_1 (Conv2D)	(None, 128, 128, 32)	12832
activation_1 (Activation)	(None, 128, 128, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 128, 128, 32)	128

In []:

```
tf.keras.utils.plot_model(model1, show_shapes=True)
```

Out[28]:



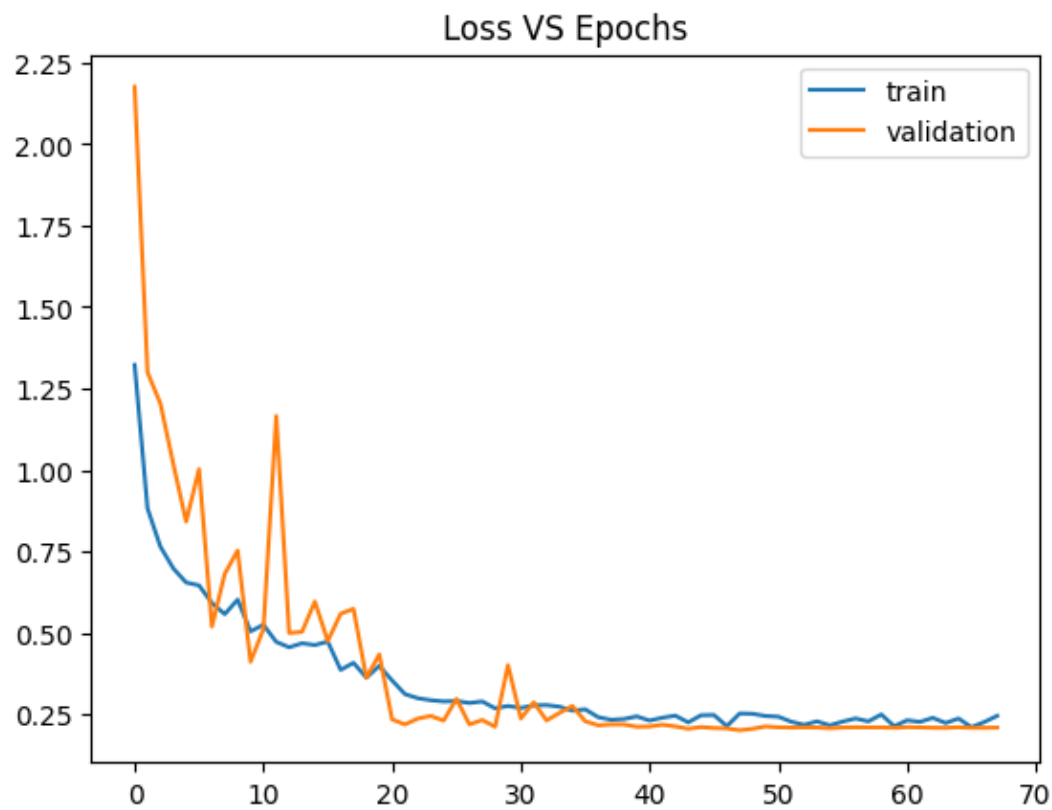
In []:

```
model_fit1 = compile_train_v1(model1, train_ds, val_ds,tb_callback,ckpt_path="/tmp/checkpoint1")
```

Epoch 1/100
100/100 [=====] - 47s 222ms/step - loss: 1.3223 - accuracy: 0.4305 - val_loss: 2.1755 - val_accuracy: 0.2912 - lr: 0.0010
Epoch 2/100
100/100 [=====] - 19s 182ms/step - loss: 0.8822 - accuracy: 0.6399 - val_loss: 1.2996 - val_accuracy: 0.4787 - lr: 0.0010
Epoch 3/100
100/100 [=====] - 19s 180ms/step - loss: 0.7648 - accuracy: 0.7181 - val_loss: 1.2037 - val_accuracy: 0.4625 - lr: 0.0010
Epoch 4/100
100/100 [=====] - 18s 175ms/step - loss: 0.6984 - accuracy: 0.7597 - val_loss: 1.0190 - val_accuracy: 0.6212 - lr: 0.0010
Epoch 5/100
100/100 [=====] - 19s 180ms/step - loss: 0.6549 - accuracy: 0.7813 - val_loss: 0.8417 - val_accuracy: 0.6363 - lr: 0.0010
Epoch 6/100
100/100 [=====] - 19s 184ms/step - loss: 0.6462 - accuracy: 0.7735 - val_loss: 1.0029 - val_accuracy: 0.6413 - lr: 0.0010
Epoch 7/100
100/100 [=====] - 19s 176ms/step - loss: 0.5921 - accuracy: 0.8021 - val_loss: 1.0029 - val_accuracy: 0.6413 - lr: 0.0010

In []:

```
epochs = model_fit1.epoch  
loss = model_fit1.history["loss"]  
val_loss = model_fit1.history["val_loss"]  
plt.plot(epochs, loss, label="train")  
plt.plot(epochs, val_loss, label="validation")  
plt.legend()  
plt.title("Loss VS Epochs")  
plt.show()
```



In []:

```
epochs = model_fit1.epoch
loss = model_fit1.history["accuracy"]
val_loss = model_fit1.history["val_accuracy"]
plt.plot(epochs, loss, label="train")
plt.plot(epochs, val_loss, label="validation")
plt.legend()
plt.title("Loss VS Epochs")
plt.show()
```



In []:

```
print_accuracy_stats(model1, test_ds, class_names, "/tmp/checkpoint1")
# Test Accuracy for CNN from scratch is 91.74%
```

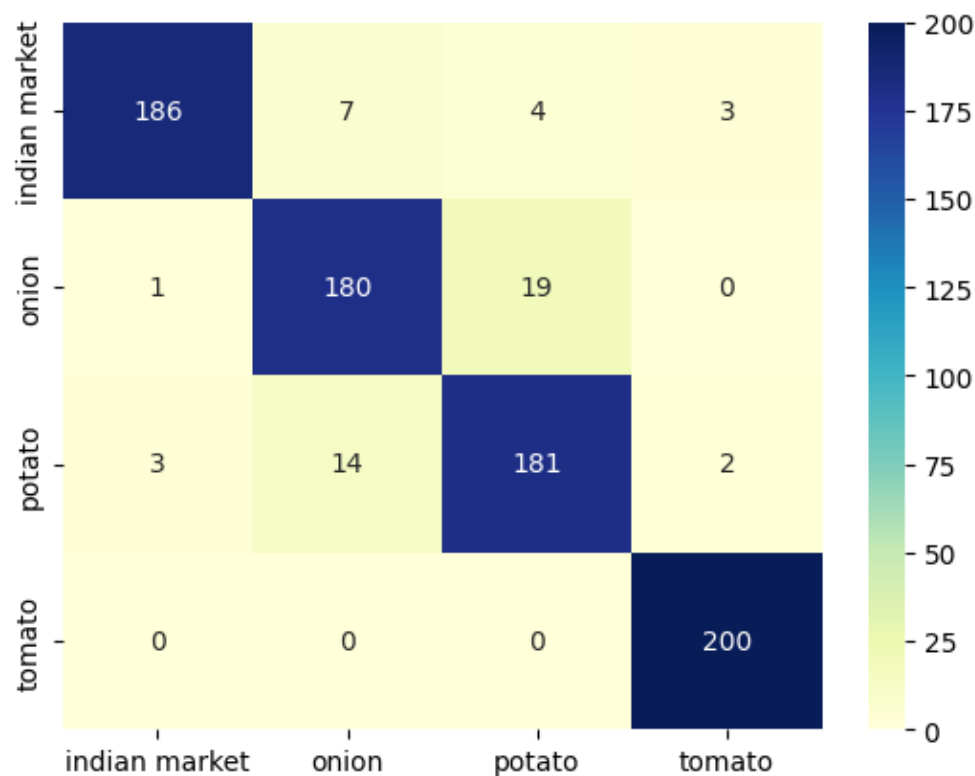
11/11 [=====] - 3s 225ms/step

Accuracy: 88.03%

In []:

```
plot_confusion_matrix(model1, val_ds, class_names, "/tmp/checkpoint1")  
# Confusion matrix for validation data
```

25/25 [=====] - 2s 86ms/step



In []:

```
print_accuracy_stats(model1, val_ds, class_names, "/tmp/checkpoint1")  
# TValidation Accuracy for CNN from scratch is 91.74%
```

25/25 [=====] - 3s 119ms/step

Accuracy: 93.38%

In []:

```
print_accuracy_stats(model1, test_ds, class_names, "/tmp/checkpoint1")
```


11/11 [=====] - 1s 107ms/step

Accuracy: 88.03%

In []:

```
plot_confusion_matrix(model1, test_ds, class_names, "/tmp/checkpoint1")
```

11/11 [=====] - 1s 108ms/step



In []:

```
#!/usr/bin/env python3
#%tensorboard --logdir logs
```

VGG

In []:

```
#!/usr/bin/env python3
#%load_ext tensorboard
log_folder='logs/2'
#%reload_ext tensorboard
#!/usr/bin/env python3
#%load_ext tensorboard
#%rm -rf logs
tb_callback=TensorBoard(log_dir=log_folder,histogram_freq=1)
```

In []:

```

pretrained_model = tf.keras.applications.VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
pretrained_model.trainable=False
vgg16_model = tf.keras.Sequential([
    pretrained_model,
    tf.keras.layers.Flatten(),
    #tf.keras.layers.Dense(10, activation='softmax')

    layers.Dense(units=512, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(units=256, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(units=128, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(units=64, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(units=32, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(units=16, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(units=8, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.1),

    layers.Dense(units=4, activation='softmax')
])

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)

58889256/58889256 [=====] - 0s 0us/step

In []:

```
vgg16_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 8, 8, 512)	14714688
flatten_1 (Flatten)	(None, 32768)	0
dense_8 (Dense)	(None, 512)	16777728
activation_13 (Activation)	(None, 512)	0
batch_normalization_13 (Batch Normalization)	(None, 512)	2048
dropout_7 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 256)	131328
activation_14 (Activation)	(None, 256)	0
batch_normalization_14 (Batch Normalization)	(None, 256)	1024
dropout_8 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 128)	32896
activation_15 (Activation)	(None, 128)	0
batch_normalization_15 (Batch Normalization)	(None, 128)	512
dropout_9 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 64)	8256
activation_16 (Activation)	(None, 64)	0
batch_normalization_16 (Batch Normalization)	(None, 64)	256
dropout_10 (Dropout)	(None, 64)	0
dense_12 (Dense)	(None, 32)	2080
activation_17 (Activation)	(None, 32)	0
batch_normalization_17 (Batch Normalization)	(None, 32)	128
dropout_11 (Dropout)	(None, 32)	0
dense_13 (Dense)	(None, 16)	528
activation_18 (Activation)	(None, 16)	0
batch_normalization_18 (Batch Normalization)	(None, 16)	64
dropout_12 (Dropout)	(None, 16)	0
dense_14 (Dense)	(None, 8)	136


```

activation_19 (Activation) (None, 8) 0

batch_normalization_19 (Batch Normalization) (None, 8) 32

dropout_13 (Dropout) (None, 8) 0

dense_15 (Dense) (None, 4) 36

```

```

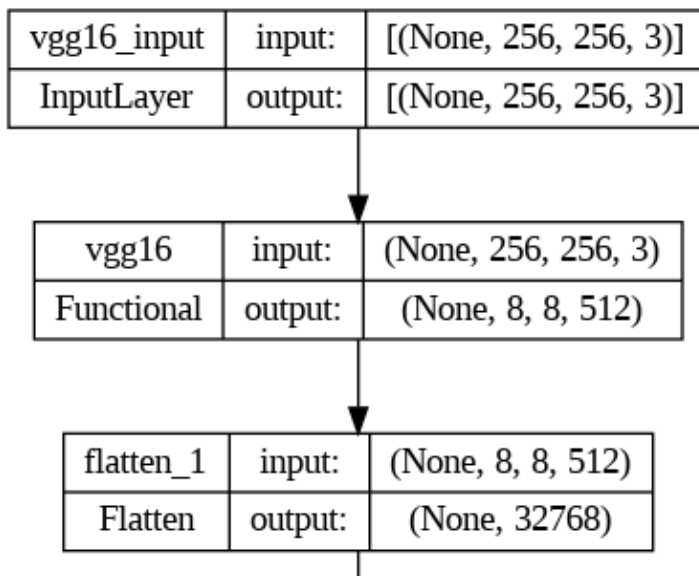
=====
Total params: 31,671,740
Trainable params: 16,955,020
Non-trainable params: 14,716,720

```

In []:

```
tf.keras.utils.plot_model(vgg16_model, show_shapes=True)
```

Out[41]:



In []:

```
model_fit2 = compile_train_v1(vgg16_model, train_ds, val_ds, tb_callback, ckpt_path="/tmp/checkpo
```

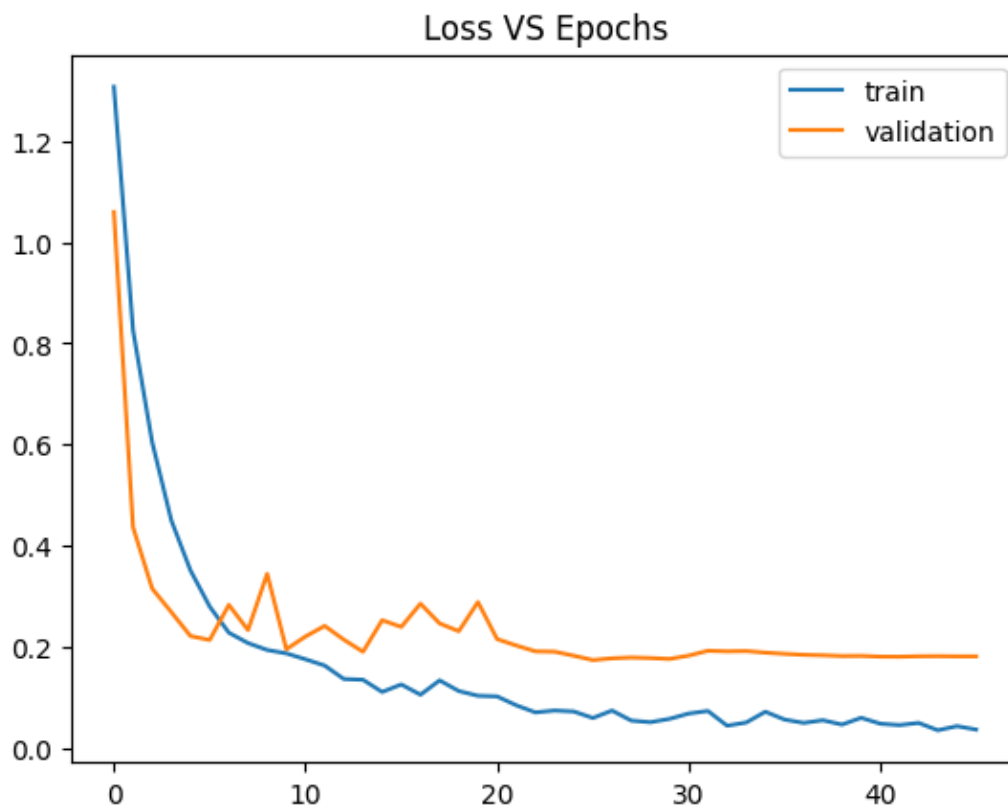
```

Epoch 1/100
100/100 [=====] - 50s 373ms/step - loss: 1.3063 - acc
uracy: 0.4399 - val_loss: 1.0592 - val_accuracy: 0.5537 - lr: 0.0010
Epoch 2/100
100/100 [=====] - 32s 315ms/step - loss: 0.8256 - acc
uracy: 0.7021 - val_loss: 0.4367 - val_accuracy: 0.8512 - lr: 0.0010
Epoch 3/100
100/100 [=====] - 31s 310ms/step - loss: 0.6044 - acc
uracy: 0.7972 - val_loss: 0.3156 - val_accuracy: 0.8975 - lr: 0.0010
Epoch 4/100
100/100 [=====] - 32s 320ms/step - loss: 0.4501 - acc
uracy: 0.8570 - val_loss: 0.2691 - val_accuracy: 0.9112 - lr: 0.0010
Epoch 5/100
100/100 [=====] - 31s 307ms/step - loss: 0.3513 - acc
uracy: 0.8936 - val_loss: 0.2220 - val_accuracy: 0.9275 - lr: 0.0010
Epoch 6/100
100/100 [=====] - 32s 315ms/step - loss: 0.2802 - acc
uracy: 0.9221 - val_loss: 0.2142 - val_accuracy: 0.9300 - lr: 0.0010
Epoch 7/100
100/100 [=====] - 32s 315ms/step - loss: 0.2402 - acc
uracy: 0.9401 - val_loss: 0.2142 - val_accuracy: 0.9300 - lr: 0.0010

```

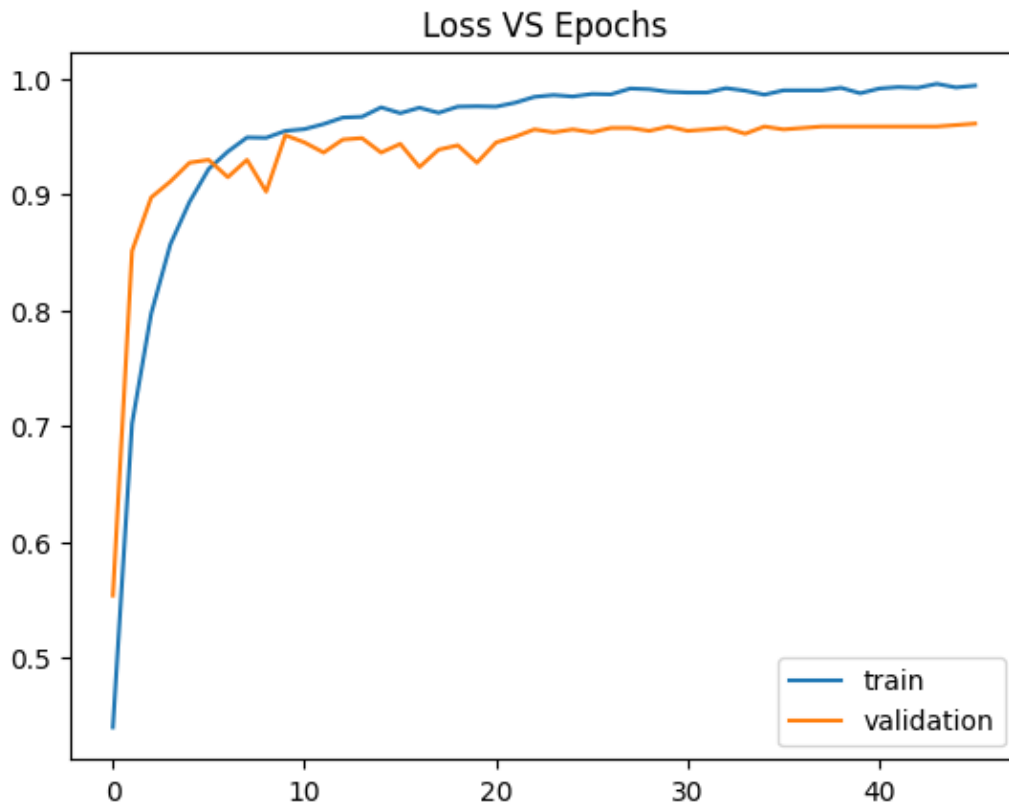
In []:

```
epochs = model_fit2.epoch  
loss = model_fit2.history["loss"]  
val_loss = model_fit2.history["val_loss"]  
plt.plot(epochs, loss, label="train")  
plt.plot(epochs, val_loss, label="validation")  
plt.legend()  
plt.title("Loss VS Epochs")  
plt.show()
```



In []:

```
epochs = model_fit2.epoch  
loss = model_fit2.history["accuracy"]  
val_loss = model_fit2.history["val_accuracy"]  
plt.plot(epochs, loss, label="train")  
plt.plot(epochs, val_loss, label="validation")  
plt.legend()  
plt.title("Loss VS Epochs")  
plt.show()
```



In []:

```
print_accuracy_stats(vgg16_model, test_ds, class_names, "/tmp/checkpoint2")
```


11/11 [=====] - 7s 708ms/step

Accuracy: 87.18%

In []:

```
plot_confusion_matrix(vgg16_model, val_ds, class_names, "/tmp/checkpoint2")
```

25/25 [=====] - 4s 181ms/step



In []:

```
print_accuracy_stats(vgg16_model, val_ds, class_names, "/tmp/checkpoint2")
```

25/25 [=====] - 4s 178ms/step

Accuracy: 96.12%

In []:

```
print_accuracy_stats(vgg16_model, test_ds, class_names, "/tmp/checkpoint2")
```

11/11 [=====] - 2s 171ms/step

Accuracy: 87.18%

In []:

```
plot_confusion_matrix(vgg16_model, test_ds, class_names, "/tmp/checkpoint2")
```

11/11 [=====] - 2s 233ms/step



In []:

```
#!/usr/bin/env python3
#%tensorboard --logdir logs
```

ResNet

In []:

```
#!/usr/bin/env python3
#%load_ext tensorboard
log_folder='logs/3'
#%reload_ext tensorboard
#!/usr/bin/env python3
#%load_ext tensorboard
log_folder='logs/3'
#%reload_ext tensorboard
#!/usr/bin/env python3
#%load_ext tensorboard
log_folder='logs/3'
#%reload_ext tensorboard
#!/usr/bin/env python3
#%load_ext tensorboard
log_folder='logs/3'
#%reload_ext tensorboard
#!/usr/bin/env python3
#%load_ext tensorboard
log_folder='logs/3'
#%reload_ext tensorboard
tb_callback=TensorBoard(log_dir=log_folder,histogram_freq=1)
```

In []:

```

pretrained_model = tf.keras.applications.resnet_v2.ResNet152V2(weights='imagenet', include_top=True)
pretrained_model.trainable=False
resnet = tf.keras.Sequential([
    pretrained_model,
    tf.keras.layers.Flatten(),
    #tf.keras.layers.Dense(10, activation='softmax')

    layers.Dense(units=128, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(units=64, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(units=32, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(units=16, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(units=8, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.1),

    layers.Dense(units=4, activation='softmax')
])

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2_weights_tf_dim_ordering_tf_kernels_notop.h5)

234545216/234545216 [=====] - 1s 0us/step

In []:

```
resnet.summary()
```


Model: "sequential_1"

Layer (type)	Output Shape	Param #
resnet152v2 (Functional)	(None, 8, 8, 2048)	58331648
flatten_2 (Flatten)	(None, 131072)	0
dense_16 (Dense)	(None, 128)	16777344
activation_20 (Activation)	(None, 128)	0
batch_normalization_20 (Batch Normalization)	(None, 128)	512
dropout_14 (Dropout)	(None, 128)	0
dense_17 (Dense)	(None, 64)	8256
activation_21 (Activation)	(None, 64)	0
batch_normalization_21 (Batch Normalization)	(None, 64)	256
dropout_15 (Dropout)	(None, 64)	0
dense_18 (Dense)	(None, 32)	2080
activation_22 (Activation)	(None, 32)	0
batch_normalization_22 (Batch Normalization)	(None, 32)	128
dropout_16 (Dropout)	(None, 32)	0
dense_19 (Dense)	(None, 16)	528
activation_23 (Activation)	(None, 16)	0
batch_normalization_23 (Batch Normalization)	(None, 16)	64
dropout_17 (Dropout)	(None, 16)	0
dense_20 (Dense)	(None, 8)	136
activation_24 (Activation)	(None, 8)	0
batch_normalization_24 (Batch Normalization)	(None, 8)	32
dropout_18 (Dropout)	(None, 8)	0
dense_21 (Dense)	(None, 4)	36
Total params: 75,121,020		
Trainable params: 16,788,876		
Non-trainable params: 58,332,144		

In []:

```
tf.keras.utils.plot_model(resnet, show_shapes=True)
```

Out[54]:



In []:

```
model_fit3 = compile_train_v1(resnet, train_ds, val_ds, tb_callback, ckpt_path="/tmp/checkpoint3")
```

Epoch 1/100
100/100 [=====] - 78s 615ms/step - loss: 0.9522 - accuracy: 0.6089 - val_loss: 0.3988 - val_accuracy: 0.9038 - lr: 0.0010

Epoch 2/100
100/100 [=====] - 60s 598ms/step - loss: 0.5206 - accuracy: 0.8326 - val_loss: 0.1724 - val_accuracy: 0.9550 - lr: 0.0010

Epoch 3/100
100/100 [=====] - 58s 571ms/step - loss: 0.3361 - accuracy: 0.9127 - val_loss: 0.1256 - val_accuracy: 0.9675 - lr: 0.0010

Epoch 4/100
100/100 [=====] - 62s 615ms/step - loss: 0.2272 - accuracy: 0.9452 - val_loss: 0.1038 - val_accuracy: 0.9725 - lr: 0.0010

Epoch 5/100
100/100 [=====] - 59s 586ms/step - loss: 0.1667 - accuracy: 0.9643 - val_loss: 0.0837 - val_accuracy: 0.9762 - lr: 0.0010

Epoch 6/100
100/100 [=====] - 56s 557ms/step - loss: 0.1320 - accuracy: 0.9712 - val_loss: 0.0715 - val_accuracy: 0.9762 - lr: 0.0010

Epoch 7/100
100/100 [=====] - 55s 547ms/step - loss: 0.0979 - accuracy: 0.9778 - val_loss: 0.0797 - val_accuracy: 0.9762 - lr: 0.0010

Epoch 8/100
100/100 [=====] - 57s 567ms/step - loss: 0.0827 - accuracy: 0.9853 - val_loss: 0.0850 - val_accuracy: 0.9787 - lr: 0.0010

Epoch 9/100
100/100 [=====] - 54s 535ms/step - loss: 0.0839 - accuracy: 0.9812 - val_loss: 0.1002 - val_accuracy: 0.9750 - lr: 0.0010

Epoch 10/100
100/100 [=====] - 58s 575ms/step - loss: 0.0655 - accuracy: 0.9872 - val_loss: 0.0962 - val_accuracy: 0.9737 - lr: 0.0010

Epoch 11/100
100/100 [=====] - 58s 573ms/step - loss: 0.0589 - accuracy: 0.9897 - val_loss: 0.1111 - val_accuracy: 0.9725 - lr: 0.0010

Epoch 12/100
100/100 [=====] - 54s 539ms/step - loss: 0.0504 - accuracy: 0.9909 - val_loss: 0.1165 - val_accuracy: 0.9725 - lr: 0.0010

Epoch 13/100
100/100 [=====] - 54s 538ms/step - loss: 0.0683 - accuracy: 0.9844 - val_loss: 0.1389 - val_accuracy: 0.9700 - lr: 0.0010

Epoch 14/100
100/100 [=====] - 54s 535ms/step - loss: 0.0550 - accuracy: 0.9900 - val_loss: 0.1019 - val_accuracy: 0.9775 - lr: 0.0010

Epoch 15/100
100/100 [=====] - 58s 575ms/step - loss: 0.0492 - accuracy: 0.9869 - val_loss: 0.1202 - val_accuracy: 0.9750 - lr: 0.0010

Epoch 16/100
100/100 [=====] - 58s 573ms/step - loss: 0.0607 - accuracy: 0.9875 - val_loss: 0.1313 - val_accuracy: 0.9737 - lr: 0.0010

Epoch 17/100
100/100 [=====] - 58s 577ms/step - loss: 0.0542 - accuracy: 0.9878 - val_loss: 0.1264 - val_accuracy: 0.9762 - lr: 0.0010

Epoch 18/100
100/100 [=====] - 57s 562ms/step - loss: 0.0476 - accuracy: 0.9903 - val_loss: 0.1093 - val_accuracy: 0.9775 - lr: 0.0010

Epoch 19/100
100/100 [=====] - 55s 551ms/step - loss: 0.0573 - accuracy: 0.9865 - val_loss: 0.1137 - val_accuracy: 0.9750 - lr: 0.0010

Epoch 20/100
100/100 [=====] - 54s 539ms/step - loss: 0.0420 - accuracy: 0.9922 - val_loss: 0.1111 - val_accuracy: 0.9787 - lr: 0.0010

Epoch 21/100
100/100 [=====] - 55s 542ms/step - loss: 0.0460 - accuracy: 0.9903 - val_loss: 0.1098 - val_accuracy: 0.9787 - lr: 1.0000e-04


Epoch 22/100

```
100/100 [=====] - 54s 539ms/step - loss: 0.0383 - accuracy: 0.9916 - val_loss: 0.1115 - val_accuracy: 0.9787 - lr: 1.0000e-04
Epoch 23/100
100/100 [=====] - 61s 603ms/step - loss: 0.0536 - accuracy: 0.9903 - val_loss: 0.1126 - val_accuracy: 0.9800 - lr: 1.0000e-04
Epoch 24/100
100/100 [=====] - 54s 538ms/step - loss: 0.0329 - accuracy: 0.9944 - val_loss: 0.1125 - val_accuracy: 0.9800 - lr: 1.0000e-04
Epoch 25/100
100/100 [=====] - 58s 578ms/step - loss: 0.0331 - accuracy: 0.9931 - val_loss: 0.1128 - val_accuracy: 0.9800 - lr: 1.0000e-04
Epoch 26/100
100/100 [=====] - 54s 536ms/step - loss: 0.0326 - accuracy: 0.9947 - val_loss: 0.1141 - val_accuracy: 0.9800 - lr: 1.0000e-04
```

In []:

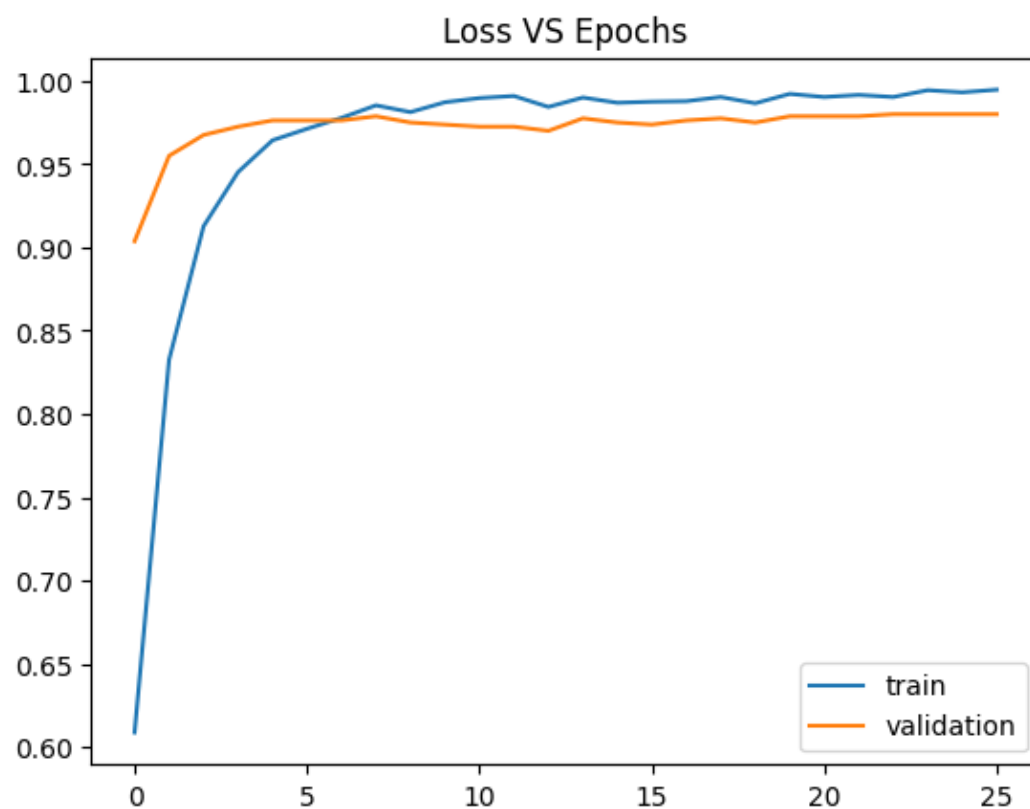
```
epochs = model_fit3.epoch
loss = model_fit3.history["loss"]
val_loss = model_fit3.history["val_loss"]
plt.plot(epochs, loss, label="train")
plt.plot(epochs, val_loss, label="validation")
plt.legend()
plt.title("Loss VS Epochs")
plt.show()
```

Loss VS Epochs



In []:

```
epochs = model_fit3.epoch  
loss = model_fit3.history["accuracy"]  
val_loss = model_fit3.history["val_accuracy"]  
plt.plot(epochs, loss, label="train")  
plt.plot(epochs, val_loss, label="validation")  
plt.legend()  
plt.title("Loss VS Epochs")  
plt.show()
```



In []:

```
print_accuracy_stats(resnet, test_ds, class_names, "/tmp/checkpoint3")
```


11/11 [=====] - 7s 446ms/step

Accuracy: 90.88%

In []:

```
plot_confusion_matrix(resnet, val_ds, class_names, "/tmp/checkpoint3")
```

25/25 [=====] - 7s 281ms/step



In []:

```
print_accuracy_stats(resnet, val_ds, class_names, "/tmp/checkpoint3")
```

25/25 [=====] - 7s 275ms/step

Accuracy: 98.00%

In []:

```
print_accuracy_stats(resnet, test_ds, class_names, "/tmp/checkpoint3")
```


11/11 [=====] - 3s 271ms/step

Accuracy: 90.88%

In []:

```
plot_confusion_matrix(resnet, test_ds, class_names, "/tmp/checkpoint3")
```

11/11 [=====] - 3s 276ms/step



In []:

```
##%tensorboard --logdir logs
```

In []:

Mobilenet

In []:

```
##%load_ext tensorboard
log_folder='logs/4'
##%reload_ext tensorboard
#!rm -rf logs
tb_callback=TensorBoard(log_dir=log_folder,histogram_freq=1)
```

In []:

```

pretrained_model = tf.keras.applications.mobilenet_v2.MobileNetV2(weights='imagenet', include_top=True)
pretrained_model.trainable=False
mobilenet = tf.keras.Sequential([
    pretrained_model,
    tf.keras.layers.Flatten(),
    #tf.keras.layers.Dense(10, activation='softmax')

    layers.Dense(units=128, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(units=64, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(units=32, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(units=16, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(units=8, kernel_regularizer = L2Reg),
    layers.Activation("relu"),
    layers.BatchNormalization(),
    layers.Dropout(0.1),

    layers.Dense(units=4, activation='softmax')
])

```

WARNING:tensorflow:`input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5
 (https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5)
 9406464/9406464 [=====] - 0s 0us/step

In []:

```
mobilenet.summary()
```

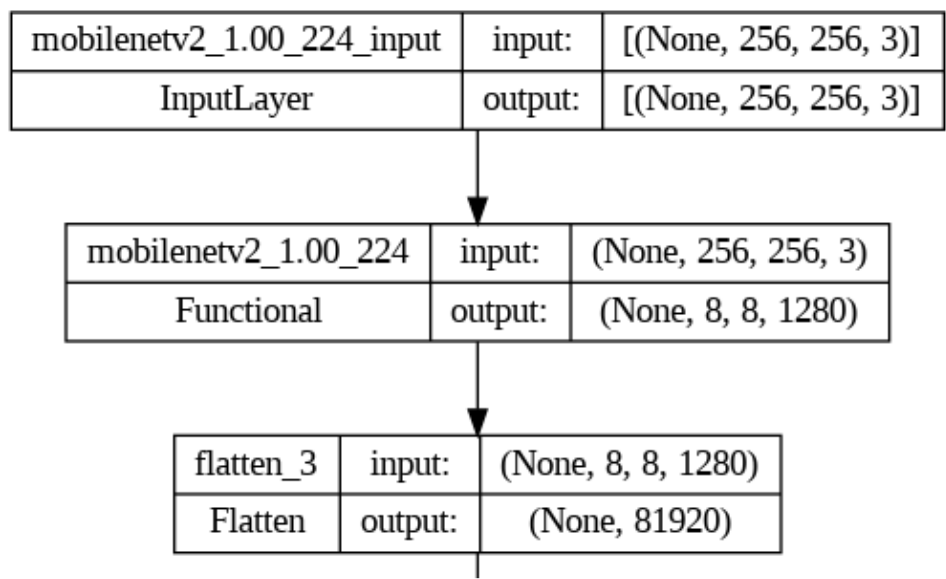
Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
mobilenetv2_1.00_224 (Functional)	(None, 8, 8, 1280)	2257984
flatten_3 (Flatten)	(None, 81920)	0
dense_22 (Dense)	(None, 128)	10485888
activation_25 (Activation)	(None, 128)	0
batch_normalization_25 (Batch Normalization)	(None, 128)	512
dropout_19 (Dropout)	(None, 128)	0
dense_23 (Dense)	(None, 64)	8256
activation_26 (Activation)	(None, 64)	0
batch_normalization_26 (Batch Normalization)	(None, 64)	256
dropout_20 (Dropout)	(None, 64)	0
dense_24 (Dense)	(None, 32)	2080
activation_27 (Activation)	(None, 32)	0
batch_normalization_27 (Batch Normalization)	(None, 32)	128
dropout_21 (Dropout)	(None, 32)	0
dense_25 (Dense)	(None, 16)	528
activation_28 (Activation)	(None, 16)	0
batch_normalization_28 (Batch Normalization)	(None, 16)	64
dropout_22 (Dropout)	(None, 16)	0
dense_26 (Dense)	(None, 8)	136
activation_29 (Activation)	(None, 8)	0
batch_normalization_29 (Batch Normalization)	(None, 8)	32
dropout_23 (Dropout)	(None, 8)	0
dense_27 (Dense)	(None, 4)	36
=====		
Total params: 12,755,900		
Trainable params: 10,497,420		
Non-trainable params: 2,258,480		

In []:

```
tf.keras.utils.plot_model(mobilenet, show_shapes=True)
```

Out[67]:



In []:

```
model_fit4 = compile_train_v1(mobilenet, train_ds, val_ds,tb_callback,ckpt_path="/tmp/checkpoint")
```

Epoch 1/100
100/100 [=====] - 27s 193ms/step - loss: 0.9483 - accuracy: 0.6133 - val_loss: 0.4741 - val_accuracy: 0.8925 - lr: 0.0010

Epoch 2/100
100/100 [=====] - 18s 172ms/step - loss: 0.5342 - accuracy: 0.8376 - val_loss: 0.2118 - val_accuracy: 0.9563 - lr: 0.0010

Epoch 3/100
100/100 [=====] - 18s 179ms/step - loss: 0.3310 - accuracy: 0.9124 - val_loss: 0.1104 - val_accuracy: 0.9750 - lr: 0.0010

Epoch 4/100
100/100 [=====] - 18s 171ms/step - loss: 0.2188 - accuracy: 0.9462 - val_loss: 0.0891 - val_accuracy: 0.9762 - lr: 0.0010

Epoch 5/100
100/100 [=====] - 18s 175ms/step - loss: 0.1464 - accuracy: 0.9725 - val_loss: 0.0692 - val_accuracy: 0.9825 - lr: 0.0010

Epoch 6/100
100/100 [=====] - 17s 165ms/step - loss: 0.1146 - accuracy: 0.9759 - val_loss: 0.0944 - val_accuracy: 0.9725 - lr: 0.0010

Epoch 7/100
100/100 [=====] - 18s 168ms/step - loss: 0.0966 - accuracy: 0.9812 - val_loss: 0.0752 - val_accuracy: 0.9775 - lr: 0.0010

Epoch 8/100
100/100 [=====] - 18s 171ms/step - loss: 0.0692 - accuracy: 0.9890 - val_loss: 0.0716 - val_accuracy: 0.9812 - lr: 0.0010

Epoch 9/100
100/100 [=====] - 17s 165ms/step - loss: 0.0546 - accuracy: 0.9890 - val_loss: 0.0723 - val_accuracy: 0.9825 - lr: 0.0010

Epoch 10/100
100/100 [=====] - 18s 174ms/step - loss: 0.0608 - accuracy: 0.9850 - val_loss: 0.0900 - val_accuracy: 0.9800 - lr: 0.0010

Epoch 11/100
100/100 [=====] - 17s 169ms/step - loss: 0.0602 - accuracy: 0.9862 - val_loss: 0.0874 - val_accuracy: 0.9800 - lr: 0.0010

Epoch 12/100
100/100 [=====] - 18s 167ms/step - loss: 0.0566 - accuracy: 0.9865 - val_loss: 0.0887 - val_accuracy: 0.9825 - lr: 0.0010

Epoch 13/100
100/100 [=====] - 17s 166ms/step - loss: 0.0605 - accuracy: 0.9847 - val_loss: 0.0790 - val_accuracy: 0.9812 - lr: 0.0010

Epoch 14/100
100/100 [=====] - 18s 175ms/step - loss: 0.0518 - accuracy: 0.9919 - val_loss: 0.0946 - val_accuracy: 0.9787 - lr: 0.0010

Epoch 15/100
100/100 [=====] - 17s 165ms/step - loss: 0.0529 - accuracy: 0.9875 - val_loss: 0.1018 - val_accuracy: 0.9775 - lr: 0.0010

Epoch 16/100
100/100 [=====] - 18s 169ms/step - loss: 0.0603 - accuracy: 0.9894 - val_loss: 0.0937 - val_accuracy: 0.9812 - lr: 0.0010

Epoch 17/100
100/100 [=====] - 17s 169ms/step - loss: 0.0376 - accuracy: 0.9931 - val_loss: 0.0915 - val_accuracy: 0.9787 - lr: 0.0010

Epoch 18/100
100/100 [=====] - 19s 184ms/step - loss: 0.0350 - accuracy: 0.9950 - val_loss: 0.1259 - val_accuracy: 0.9775 - lr: 0.0010

Epoch 19/100
100/100 [=====] - 17s 167ms/step - loss: 0.0397 - accuracy: 0.9916 - val_loss: 0.1273 - val_accuracy: 0.9750 - lr: 0.0010

Epoch 20/100
100/100 [=====] - 17s 165ms/step - loss: 0.0516 - accuracy: 0.9887 - val_loss: 0.1298 - val_accuracy: 0.9762 - lr: 0.0010


Epoch 21/100
100/100 [=====] - 18s 170ms/step - loss: 0.0481 - accuracy: 0.9900 - val_loss: 0.1184 - val_accuracy: 0.9787 - lr: 1.0000e-04

Epoch 22/100

```
100/100 [=====] - 17s 169ms/step - loss: 0.0321 - accuracy: 0.9937 - val_loss: 0.1191 - val_accuracy: 0.9775 - lr: 1.0000e-04  
Epoch 23/100  
100/100 [=====] - 17s 163ms/step - loss: 0.0303 - accuracy: 0.9937 - val_loss: 0.1208 - val_accuracy: 0.9775 - lr: 1.0000e-04  
Epoch 24/100  
100/100 [=====] - 18s 175ms/step - loss: 0.0295 - accuracy: 0.9947 - val_loss: 0.1208 - val_accuracy: 0.9762 - lr: 1.0000e-04  
Epoch 25/100  
100/100 [=====] - 18s 171ms/step - loss: 0.0196 - accuracy: 0.9975 - val_loss: 0.1206 - val_accuracy: 0.9762 - lr: 1.0000e-04
```


In []:

```
epochs = model_fit4.epoch  
loss = model_fit4.history["loss"]  
val_loss = model_fit4.history["val_loss"]  
plt.plot(epochs, loss, label="train")  
plt.plot(epochs, val_loss, label="validation")  
plt.legend()  
plt.title("Loss VS Epochs")  
plt.show()
```



In []:

```
epochs = model_fit4.epoch  
loss = model_fit4.history["accuracy"]  
val_loss = model_fit4.history["val_accuracy"]  
plt.plot(epochs, loss, label="train")  
plt.plot(epochs, val_loss, label="validation")  
plt.legend()  
plt.title("Loss VS Epochs")  
plt.show()
```



In []:

```
print_accuracy_stats(mobilenet, test_ds, class_names, "/tmp/checkpoint4")
```

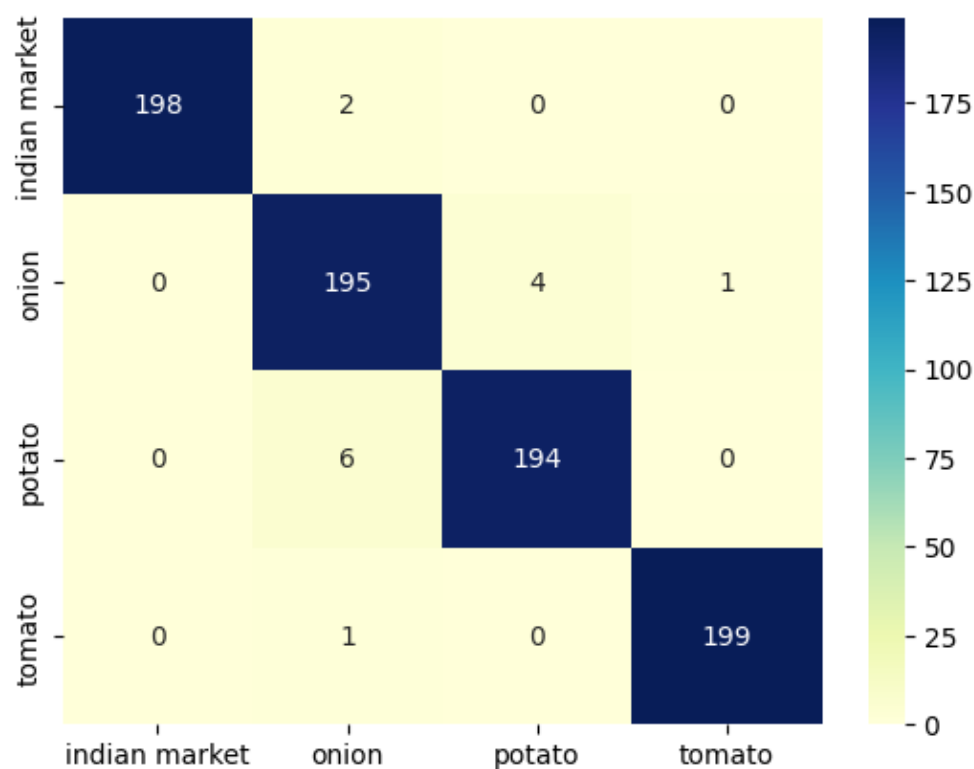
11/11 [=====] - 4s 251ms/step

Accuracy: 92.31%

In []:

```
plot_confusion_matrix(mobilenet, val_ds, class_names, "/tmp/checkpoint4")
```

25/25 [=====] - 2s 94ms/step



In []:

```
print_accuracy_stats(mobilenet, val_ds, class_names, "/tmp/checkpoint4")
```

25/25 [=====] - 3s 115ms/step

Accuracy: 98.25%

In []:

```
print_accuracy_stats(mobilenet, test_ds, class_names, "/tmp/checkpoint4")
```


11/11 [=====] - 1s 111ms/step

Accuracy: 92.31%

In []:

```
plot_confusion_matrix(mobilenet, test_ds, class_names, "/tmp/checkpoint4")
```

11/11 [=====] - 1s 114ms/step



In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

```
%tensorboard --logdir logs
```

<IPython.core.display.Javascript object>

In []:

In []:

In []: