

Problem Statement : Delivery time estimation using Artificial Neural Networks

In [1]:

```
import pandas as pd
import numpy as np
import datetime as dt
import category_encoders as ce
import numpy as np
import pandas as pd
import regex as re
import datetime as dt
import seaborn as sns
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = (20, 6)
import warnings
warnings.filterwarnings("ignore")
from sklearn.neighbors import LocalOutlierFactor
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
```

In []:

In [2]:

```
!gdown 1jq16ByYD4HHd9ESL8_v5PKn33PGHjE2k
```

```
/bin/bash: gdown: command not found
```

In [3]:

```
pip install category_encoders
```

Requirement already satisfied: category_encoders in /opt/conda/lib/python3.7/site-packages (2.6.0)
Requirement already satisfied: scipy>=1.0.0 in /opt/conda/lib/python3.7/site-packages (from category_encoders) (1.7.3)
Requirement already satisfied: numpy>=1.14.0 in /opt/conda/lib/python3.7/site-packages (from category_encoders) (1.21.6)
Requirement already satisfied: patsy>=0.5.1 in /opt/conda/lib/python3.7/site-packages (from category_encoders) (0.5.3)
Requirement already satisfied: statsmodels>=0.9.0 in /opt/conda/lib/python3.7/site-packages (from category_encoders) (0.13.5)
Requirement already satisfied: scikit-learn>=0.20.0 in /opt/conda/lib/python3.7/site-packages (from category_encoders) (1.0.2)
Requirement already satisfied: pandas>=1.0.5 in /opt/conda/lib/python3.7/site-packages (from category_encoders) (1.3.5)
Requirement already satisfied: python-dateutil>=2.7.3 in /opt/conda/lib/python3.7/site-packages (from pandas>=1.0.5->category_encoders) (2.8.2)
Requirement already satisfied: pytz>=2017.3 in /opt/conda/lib/python3.7/site-packages (from pandas>=1.0.5->category_encoders) (2022.7.1)
Requirement already satisfied: six in /opt/conda/lib/python3.7/site-packages (from patsy>=0.5.1->category_encoders) (1.16.0)
Requirement already satisfied: joblib>=0.11 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.20.0->category_encoders) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.20.0->category_encoders) (3.1.0)
Requirement already satisfied: packaging>=21.3 in /opt/conda/lib/python3.7/site-packages (from statsmodels>=0.9.0->category_encoders) (23.0)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: <https://pip.pypa.io/warnings/venv> (<https://pip.pypa.io/warnings/venv>)
WARNING: There was an error checking the latest version of pip.
Note: you may need to restart the kernel to use updated packages.

In [4]:

```
pip install seaborn
```

Requirement already satisfied: seaborn in /opt/conda/lib/python3.7/site-packages (0.12.2)
Requirement already satisfied: typing_extensions in /opt/conda/lib/python3.7/site-packages (from seaborn) (4.4.0)
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in /opt/conda/lib/python3.7/site-packages (from seaborn) (3.5.3)
Requirement already satisfied: numpy!=1.24.0,>=1.17 in /opt/conda/lib/python3.7/site-packages (from seaborn) (1.21.6)
Requirement already satisfied: pandas>=0.25 in /opt/conda/lib/python3.7/site-packages (from seaborn) (1.3.5)
Requirement already satisfied: pillow>=6.2.0 in /opt/conda/lib/python3.7/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (9.4.0)
Requirement already satisfied: packaging>=20.0 in /opt/conda/lib/python3.7/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (23.0)
Requirement already satisfied: python-dateutil>=2.7 in /opt/conda/lib/python3.7/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.8.2)
Requirement already satisfied: fonttools>=4.22.0 in /opt/conda/lib/python3.7/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.38.0)
Requirement already satisfied: pyparsing>=2.2.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.0.9)
Requirement already satisfied: kiwisolver>=1.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.4)
Requirement already satisfied: cycler>=0.10 in /opt/conda/lib/python3.7/site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.11.0)
Requirement already satisfied: pytz>=2017.3 in /opt/conda/lib/python3.7/site-packages (from pandas>=0.25->seaborn) (2022.7.1)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.7/site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.16.0)
WARNING: Running pip as the 'root' user can result in broken permissions a
nd conflicting behaviour with the system package manager. It is recommend
ed to use a virtual environment instead: <https://pip.pypa.io/warnings/venv>
(<https://pip.pypa.io/warnings/venv>)
WARNING: There was an error checking the latest version of pip.
Note: you may need to restart the kernel to use updated packages.

Reading Data

In [5]:

```
import os  
os.listdir('/kaggle/input')
```

Out[5]:

```
['newdataset']
```

In [2]:

```
df=pd.read_csv('dataset.csv')
```


In [3]:

df

Out[3]:

	market_id	created_at	actual_delivery_time	store_id	store
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	
...
197423	1.0	2015-02-17 00:19:41	2015-02-17 01:24:48	a914ecef9c12ffdb9bede64bb703d877	
197424	1.0	2015-02-13 00:01:59	2015-02-13 00:58:22	a914ecef9c12ffdb9bede64bb703d877	
197425	1.0	2015-01-24 04:46:08	2015-01-24 05:36:16	a914ecef9c12ffdb9bede64bb703d877	
197426	1.0	2015-02-01 18:18:15	2015-02-01 19:23:22	c81e155d85dae5430a8cee6f2242e82c	
197427	1.0	2015-02-08 19:24:33	2015-02-08 20:01:41	c81e155d85dae5430a8cee6f2242e82c	

197428 rows × 14 columns



Checking the structure & characteristics of the dataset

In [4]:

```
df.describe()
```

Out[4]:

	market_id	order_protocol	total_items	subtotal	num_distinct_items	mir
count	196441.000000	196433.000000	197428.000000	197428.000000	197428.000000	19
mean	2.978706	2.882352	3.196391	2682.331402	2.670791	
std	1.524867	1.503771	2.666546	1823.093688	1.630255	
min	1.000000	1.000000	1.000000	0.000000	1.000000	
25%	2.000000	1.000000	2.000000	1400.000000	1.000000	
50%	3.000000	3.000000	3.000000	2200.000000	2.000000	
75%	4.000000	4.000000	4.000000	3395.000000	3.000000	
max	6.000000	7.000000	411.000000	27100.000000	20.000000	1

In [5]:

```
df.describe(include='object')
```

Out[5]:

	created_at	actual_delivery_time	store_id	store_primary_ca
count	197428	197421	197428	1
unique	180985	178110	6743	
top	2015-02-11 19:50:43	2015-02-11 20:40:45	d43ab110ab2489d6b9b2caa394bf920f	an
freq	6	5	937	

In [6]:

```
df.isnull().sum()
# There are null values in couple of columns
```

Out[6]:

```
market_id          987
created_at         0
actual_delivery_time 7
store_id           0
store_primary_category 4760
order_protocol     995
total_items        0
subtotal           0
num_distinct_items 0
min_item_price     0
max_item_price     0
total_onshift_partners 16262
total_busy_partners 16262
total_outstanding_orders 16262
dtype: int64
```

In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   market_id                            196441 non-null  float64
 1   created_at                           197428 non-null  object
 2   actual_delivery_time                 197421 non-null  object
 3   store_id                             197428 non-null  object
 4   store_primary_category               192668 non-null  object
 5   order_protocol                       196433 non-null  float64
 6   total_items                          197428 non-null  int64
 7   subtotal                             197428 non-null  int64
 8   num_distinct_items                  197428 non-null  int64
 9   min_item_price                      197428 non-null  int64
10  max_item_price                      197428 non-null  int64
11  total_onshift_partners               181166 non-null  float64
12  total_busy_partners                  181166 non-null  float64
13  total_outstanding_orders             181166 non-null  float64
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB
```

Data preprocessing

In [8]:

```
df['market_id'].value_counts(normalize=True)
# 28.02 % values are from market_id 2
```

Out[8]:

```
2.0    0.280278
4.0    0.242307
1.0    0.193631
3.0    0.118595
5.0    0.091631
6.0    0.073559
Name: market_id, dtype: float64
```

In [9]:

```
df['store_id'].value_counts(normalize=True)
```

Out[9]:

```
d43ab110ab2489d6b9b2caa394bf920f    0.004746
757b505cfd34c64c85ca5b5690ee5293    0.004371
faacbcd5bf1d018912c116bf2783e9a1    0.004128
cfecdb276f634854f3ef915e2e980c31    0.003875
45c48cce2e2d7fbdea1afc51c7c6ad26    0.003652
...
adad0f2b196a1ed3e3b9d9025c397132    0.000005
2e6d9c6052e99fcdfa61d9b9da273ca2    0.000005
25daeb9b3072e9c53f66a2196a92a011    0.000005
55285adfd78a019a3245917649e29b3c    0.000005
df263d996281d984952c07998dc54358    0.000005
Name: store_id, Length: 6743, dtype: float64
```

In [10]:

```
df['store_primary_category'].value_counts(normalize=True)
# There are 74 Unique values of store_primary_category
```

Out[10]:

```
american    0.100686
pizza        0.089901
mexican      0.088749
burger       0.056875
sandwich     0.052214
...
lebanese     0.000047
belgian      0.000010
indonesian   0.000010
chocolate    0.000005
alcohol-plus-food 0.000005
Name: store_primary_category, Length: 74, dtype: float64
```

In [11]:

```
df['order_protocol'].value_counts(normalize=True)
# 27.85% of the order protocols are from 1 followed by 3 and 5
```


Out[11]:

```
1.0    0.278594
3.0    0.270825
5.0    0.225471
2.0    0.122444
4.0    0.098527
6.0    0.004042
7.0    0.000097
Name: order_protocol, dtype: float64
```

Data Visualization and cleaning


In [12]:

```
sns.countplot(data=df, x='market_id', order=df['market_id'].value_counts().index)
plt.xticks(rotation = 90)
plt.show()
# Count plot below shows number of market_id's
```




In [13]:

```
sns.countplot(data=df, x='market_id', order=df['market_id'].value_counts().index, hue='order_protocol')
plt.xticks(rotation = 90)
plt.show()
# order_protocol of 3 is highest with market_id 2 and 4
```



In [14]:

```
sns.countplot(data=df, x='store_primary_category', order=df['store_primary_category'].value_counts().index)
plt.xticks(rotation = 90)
plt.show()
# American, Pizza and Mexican are the top most items in the store_primary_category
```

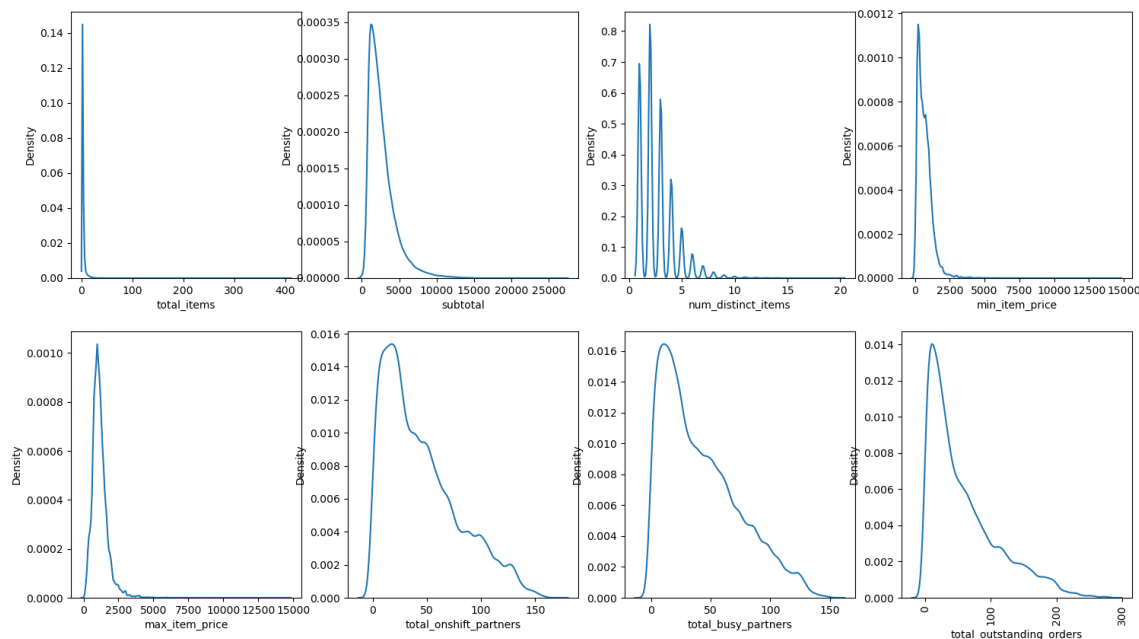


In [15]:

```
fig, axes = plt.subplots(2, 4, figsize=(18, 10))
plt.xticks(rotation = 90)
fig.suptitle('kde plot')
sns.kdeplot(ax=axes[0, 0], data=df, x='total_items')
sns.kdeplot(ax=axes[0, 1], data=df, x='subtotal')
sns.kdeplot(ax=axes[0, 2], data=df, x='num_distinct_items')
sns.kdeplot(ax=axes[0, 3], data=df, x='min_item_price')
sns.kdeplot(ax=axes[1, 0], data=df, x='max_item_price')
sns.kdeplot(ax=axes[1, 1], data=df, x='total_onshift_partners')
sns.kdeplot(ax=axes[1, 2], data=df, x='total_busy_partners')
sns.kdeplot(ax=axes[1, 3], data=df, x='total_outstanding_orders')
plt.show()
```

```
# kde plot for all the continuous variables
```

kde plot



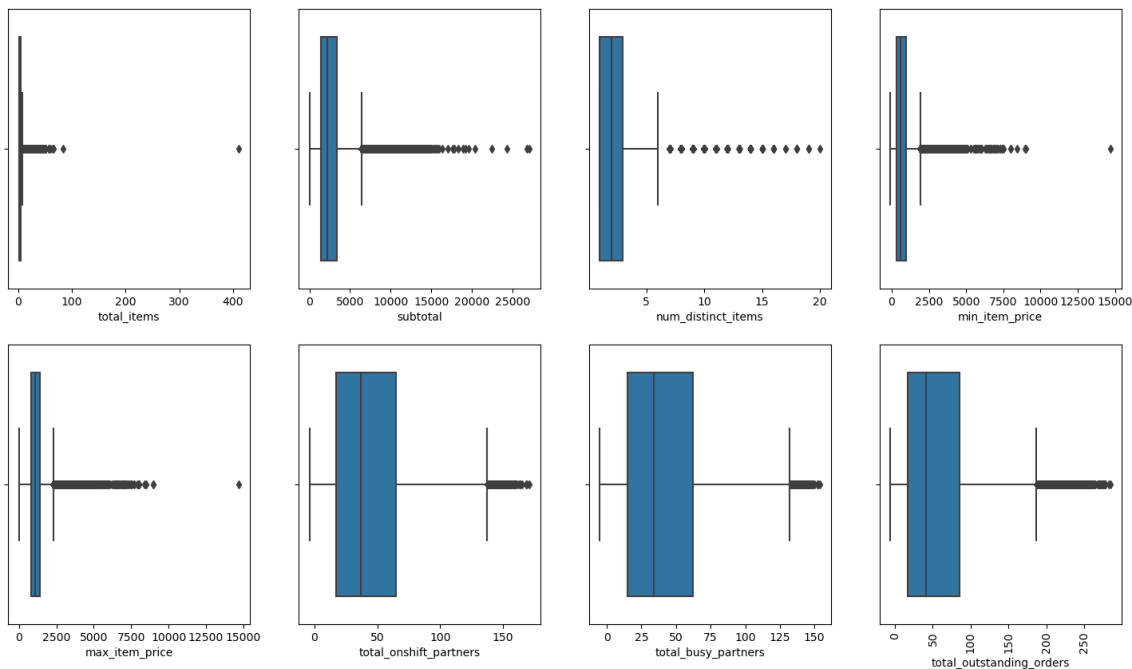
In [16]:

```

fig, axes = plt.subplots(2, 4, figsize=(18, 10))
plt.xticks(rotation = 90)
fig.suptitle('Box plot')
sns.boxplot(ax=axes[0, 0], data=df, x='total_items')
sns.boxplot(ax=axes[0, 1], data=df, x='subtotal')
sns.boxplot(ax=axes[0, 2], data=df, x='num_distinct_items' )
sns.boxplot(ax=axes[0, 3], data=df, x='min_item_price' )
sns.boxplot(ax=axes[1, 0], data=df, x='max_item_price')
sns.boxplot(ax=axes[1, 1], data=df, x='total_onshift_partners')
sns.boxplot(ax=axes[1, 2], data=df, x='total_busy_partners')
sns.boxplot(ax=axes[1, 3], data=df, x='total_outstanding_orders')
plt.show()
# Boxplot for all the continuous variables
# There are outliers in each variable

```

Box plot




In [17]:

```
fig, ax = plt.subplots(figsize=(15,10))
sns.heatmap(df.corr(method = 'pearson'),annot=True,ax=ax)
# Pearson corr
```

Out[17]:

<AxesSubplot:>




In [18]:

```
fig, ax = plt.subplots(figsize=(15,10))
sns.heatmap(df.corr(method='spearman'),annot=True,ax=ax)
```

Out[18]:

<AxesSubplot:>



Missing values Imputation

In [19]:

```
dff=df[['store_id','store_primary_category']]
df1=pd.DataFrame(dff.groupby(['store_id','store_primary_category'])['store_primary_category'].count())
df1 = df1.rename(columns={'store_primary_category': 'count'})
df1 = df1.reset_index()
df1=df1.sort_values(by = ['store_id','store_primary_category','count'], ascending = [True,True,False])
df1=df1.groupby(['store_id']).first().reset_index()
df=df.merge(df1, how='left', on=['store_id'])
for i in range(len(df)):
    if pd.isnull(df.iloc[i,4]):
        if pd.isnull(df.iloc[i,14])==False:
            df.iloc[i,4]=df.iloc[i,14]
df.drop(['store_primary_category_y','count'],axis=1,inplace=True)
df.rename(columns={'store_primary_category_x':'store_primary_category'},inplace=True)

# Filling missing store_primary_category with count of max occurrence based on store_id
```

In [20]:

```
df.isnull().sum()
```

Out[20]:

```
market_id          987
created_at         0
actual_delivery_time 7
store_id           0
store_primary_category 867
order_protocol     995
total_items        0
subtotal           0
num_distinct_items 0
min_item_price     0
max_item_price     0
total_onshift_partners 16262
total_busy_partners 16262
total_outstanding_orders 16262
dtype: int64
```

In [21]:

```
dff=df[['store_id','market_id']]
df1=pd.DataFrame(dff.groupby(['store_id','market_id'])['market_id'].agg('count'))
if 'market_id' in df1.columns:
    df1 = df1.rename(columns={'market_id': 'count'})
df1 = df1.reset_index()
df1=df1.sort_values(by = ['store_id','market_id','count'], ascending = [True,True,False])
df1=df1.groupby(['store_id']).first().reset_index()
df=df.merge(df1, how='left', on=['store_id'])
for i in range(len(df)):
    if pd.isnull(df.iloc[i,0]):
        if pd.isnull(df.iloc[i,14])==False:
            df.iloc[i,0]=df.iloc[i,14]
df.drop(['market_id_y','count'],axis=1,inplace=True)
df.columns=['market_id', 'created_at', 'actual_delivery_time', 'store_id',
            'store_primary_category', 'order_protocol', 'total_items', 'subtotal',
            'num_distinct_items', 'min_item_price', 'max_item_price',
            'total_onshift_partners', 'total_busy_partners',
            'total_outstanding_orders']
df.dropna(subset=['market_id', 'actual_delivery_time'],inplace=True)

# Filling missing market_id
```

In [22]:

```
df.isnull().sum()
```

Out[22]:

```
market_id          0
created_at         0
actual_delivery_time 0
store_id           0
store_primary_category 866
order_protocol     992
total_items        0
subtotal           0
num_distinct_items 0
min_item_price     0
max_item_price     0
total_onshift_partners 16261
total_busy_partners 16261
total_outstanding_orders 16261
dtype: int64
```

In [23]:

```
dff=df[['store_id','order_protocol']]
df1=pd.DataFrame(dff.groupby(['store_id','order_protocol'])['order_protocol'].agg('count'))
if 'order_protocol' in df1.columns:
    df1 = df1.rename(columns={'order_protocol': 'count'})
df1 = df1.reset_index()
df1=df1.sort_values(by = ['store_id','order_protocol','count'], ascending = [True,True,False])
df1=df1.groupby(['store_id']).first().reset_index()
df=df.merge(df1, how='left', on=['store_id'])
for i in range(len(df)):
    if pd.isnull(df.iloc[i,5]):
        if pd.isnull(df.iloc[i,14])==False:
            df.iloc[i,5]=df.iloc[i,14]
df.drop(['order_protocol_y', 'count'],axis=1,inplace=True)
df.columns=['market_id', 'created_at', 'actual_delivery_time', 'store_id',
            'store_primary_category', 'order_protocol', 'total_items', 'subtotal',
            'num_distinct_items', 'min_item_price', 'max_item_price',
            'total_onshift_partners', 'total_busy_partners',
            'total_outstanding_orders']
# filling missing market_id
```

In [24]:

```
df.isnull().sum()
```

Out[24]:

```
market_id          0
created_at         0
actual_delivery_time 0
store_id           0
store_primary_category 866
order_protocol     0
total_items        0
subtotal           0
num_distinct_items 0
min_item_price     0
max_item_price     0
total_onshift_partners 16261
total_busy_partners 16261
total_outstanding_orders 16261
dtype: int64
```

In [25]:

```
dff=df[['market_id','store_primary_category']]
df1=pd.DataFrame(dff.groupby(['market_id','store_primary_category'])['store_primary_category'].count().reset_index())
if 'store_primary_category' in df1.columns:
    df1 = df1.rename(columns={'store_primary_category': 'count'})
df1 = df1.reset_index()
df1=df1.sort_values(by = ['market_id','store_primary_category','count'], ascending = [True,True,True])
df1=df1.groupby(['market_id']).first().reset_index()
df=df.merge(df1, how='left', on=['market_id'])
for i in range(len(df)):
    if pd.isnull(df.iloc[i,4]):
        if pd.isnull(df.iloc[i,14])==False:
            df.iloc[i,4]=df.iloc[i,14]
df.drop(['store_primary_category_y','count'],axis=1,inplace=True)
df.columns=['market_id', 'created_at', 'actual_delivery_time', 'store_id',
            'store_primary_category', 'order_protocol', 'total_items', 'subtotal',
            'num_distinct_items', 'min_item_price', 'max_item_price',
            'total_onshift_partners', 'total_busy_partners',
            'total_outstanding_orders']
# filling missing store_primary_category
```

In [26]:

```
df.isnull().sum()
```

Out[26]:

```
market_id                0
created_at               0
actual_delivery_time     0
store_id                0
store_primary_category   0
order_protocol           0
total_items             0
subtotal                0
num_distinct_items       0
min_item_price           0
max_item_price           0
total_onshift_partners   16261
total_busy_partners      16261
total_outstanding_orders 16261
dtype: int64
```

Feature Engineering

In [27]:

```
df['created_at']=pd.to_datetime(df['created_at'])
df['actual_delivery_time']=pd.to_datetime(df['actual_delivery_time'])
```

In [28]:

```
df['created_at_hour']=df['created_at'].dt.hour
df['created_at_date']=df['created_at'].dt.date
```

In [29]:

```
df.isnull().sum()
```

Out[29]:

```
market_id                0
created_at               0
actual_delivery_time     0
store_id                0
store_primary_category   0
order_protocol           0
total_items             0
subtotal                0
num_distinct_items       0
min_item_price           0
max_item_price           0
total_onshift_partners   16261
total_busy_partners      16261
total_outstanding_orders 16261
created_at_hour          0
created_at_date           0
dtype: int64
```


In [30]:

```

dff=df[['market_id', 'created_at_hour',
        'store_primary_category', 'total_onshift_partners']]
df1=pd.DataFrame(dff.groupby(['market_id', 'created_at_hour',
        'store_primary_category'])['total_onshift_partners'].agg('mean'))
if 'total_onshift_partners' in df1.columns:
    df1 = df1.rename(columns={'total_onshift_partners': 'count'})
df1 = df1.reset_index()

df=df.merge(df1, how='left', on=['market_id', 'created_at_hour',
        'store_primary_category'])
for i in range(len(df)):
    if pd.isnull(df.iloc[i,11]):
        if pd.isnull(df.iloc[i,16])==False:
            df.iloc[i,11]=df.iloc[i,16]
df.drop('count',axis=1,inplace=True)
dff=df[['market_id', 'created_at_hour',
        'store_primary_category', 'total_busy_partners']]
df1=pd.DataFrame(dff.groupby(['market_id', 'created_at_hour',
        'store_primary_category'])['total_busy_partners'].agg('mean'))
if 'total_busy_partners' in df1.columns:
    df1 = df1.rename(columns={'total_busy_partners': 'count'})
df1 = df1.reset_index()

df=df.merge(df1, how='left', on=['market_id', 'created_at_hour',
        'store_primary_category'])
for i in range(len(df)):
    if pd.isnull(df.iloc[i,12]):
        if pd.isnull(df.iloc[i,16])==False:
            df.iloc[i,12]=df.iloc[i,16]
df.drop('count',axis=1,inplace=True)
dff=df[['market_id', 'created_at_hour',
        'store_primary_category', 'total_outstanding_orders']]
df1=pd.DataFrame(dff.groupby(['market_id', 'created_at_hour',
        'store_primary_category'])['total_outstanding_orders'].agg('mean'))
if 'total_outstanding_orders' in df1.columns:
    df1 = df1.rename(columns={'total_outstanding_orders': 'count'})
df1 = df1.reset_index()

df=df.merge(df1, how='left', on=['market_id', 'created_at_hour',
        'store_primary_category'])
for i in range(len(df)):
    if pd.isnull(df.iloc[i,13]):
        if pd.isnull(df.iloc[i,16])==False:
            df.iloc[i,13]=df.iloc[i,16]
df.drop('count',axis=1,inplace=True)

# Filling missing values of total_onshift_partners,total_busy_partners and total_outstan

```

In [31]:

```
df.isnull().sum()  
# There are still some missing values
```

Out[31]:

```
market_id          0  
created_at         0  
actual_delivery_time 0  
store_id           0  
store_primary_category 0  
order_protocol     0  
total_items        0  
subtotal           0  
num_distinct_items 0  
min_item_price     0  
max_item_price     0  
total_onshift_partners 3121  
total_busy_partners 3121  
total_outstanding_orders 3121  
created_at_hour    0  
created_at_date    0  
dtype: int64
```

In [32]:

```
3121/len(df)  
# We are dropping 1.58% of values
```

Out[32]:

```
0.015809095421896686
```

In [33]:

```
df.dropna(subset=['total_onshift_partners', 'total_busy_partners', 'total_outstanding_ord
```

In [34]:

```
df.isnull().sum()
# There are no null values
```

Out[34]:

```
market_id          0
created_at         0
actual_delivery_time 0
store_id           0
store_primary_category 0
order_protocol     0
total_items        0
subtotal           0
num_distinct_items 0
min_item_price     0
max_item_price     0
total_onshift_partners 0
total_busy_partners 0
total_outstanding_orders 0
created_at_hour    0
created_at_date    0
dtype: int64
```

In [35]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 194297 entries, 0 to 197417
Data columns (total 16 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   market_id                            194297 non-null float64
 1   created_at                           194297 non-null datetime64[ns]
 2   actual_delivery_time                 194297 non-null datetime64[ns]
 3   store_id                             194297 non-null object
 4   store_primary_category               194297 non-null object
 5   order_protocol                       194297 non-null float64
 6   total_items                          194297 non-null int64
 7   subtotal                             194297 non-null int64
 8   num_distinct_items                  194297 non-null int64
 9   min_item_price                      194297 non-null int64
10   max_item_price                      194297 non-null int64
11   total_onshift_partners               194297 non-null float64
12   total_busy_partners                  194297 non-null float64
13   total_outstanding_orders             194297 non-null float64
14   created_at_hour                      194297 non-null int64
15   created_at_date                      194297 non-null object
dtypes: datetime64[ns](2), float64(5), int64(6), object(3)
memory usage: 25.2+ MB
```

Delivery time in minutes

In [36]:

```
df['delivery_time_min']=((df['actual_delivery_time']-df['created_at'])/np.timedelta64(1,
```

In [37]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 194297 entries, 0 to 197417
Data columns (total 17 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   market_id                            194297 non-null float64
 1   created_at                           194297 non-null datetime64[ns]
 2   actual_delivery_time                 194297 non-null datetime64[ns]
 3   store_id                             194297 non-null object
 4   store_primary_category              194297 non-null object
 5   order_protocol                      194297 non-null float64
 6   total_items                         194297 non-null int64
 7   subtotal                           194297 non-null int64
 8   num_distinct_items                  194297 non-null int64
 9   min_item_price                      194297 non-null int64
10   max_item_price                      194297 non-null int64
11   total_onshift_partners              194297 non-null float64
12   total_busy_partners                 194297 non-null float64
13   total_outstanding_orders           194297 non-null float64
14   created_at_hour                     194297 non-null int64
15   created_at_date                     194297 non-null object
16   delivery_time_min                   194297 non-null float64
dtypes: datetime64[ns](2), float64(6), int64(6), object(3)
memory usage: 26.7+ MB
```

In [38]:

```
df['store_id'].nunique()
```

Out[38]:

6654

In [39]:

```
df['store_primary_category'].nunique()
```

Out[39]:

73

Encoding categorical Columns

In [40]:

```
encoder = ce.TargetEncoder(cols=['store_id'])
df_encoded = encoder.fit_transform(df['store_id'],df['delivery_time_min'])
df['store_id']=df_encoded
encoder = ce.TargetEncoder(cols=['store_primary_category'])
df_encoded = encoder.fit_transform(df['store_primary_category'],df['delivery_time_min'])
df['store_primary_category']=df_encoded
```

In [41]:

```
df['total_onshift_partners']=df['total_onshift_partners'].apply(lambda x : int(x))
df['total_busy_partners']=df['total_busy_partners'].apply(lambda x : int(x))
df['total_outstanding_orders']=df['total_outstanding_orders'].apply(lambda x : int(x))

# Converting to int values for all the float values as we had taken mean earlier
```

In [42]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 194297 entries, 0 to 197417
Data columns (total 17 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   market_id                            194297 non-null  float64
 1   created_at                           194297 non-null  datetime64[ns]
 2   actual_delivery_time                 194297 non-null  datetime64[ns]
 3   store_id                             194297 non-null  float64
 4   store_primary_category               194297 non-null  float64
 5   order_protocol                       194297 non-null  float64
 6   total_items                          194297 non-null  int64
 7   subtotal                             194297 non-null  int64
 8   num_distinct_items                  194297 non-null  int64
 9   min_item_price                      194297 non-null  int64
10   max_item_price                      194297 non-null  int64
11   total_onshift_partners               194297 non-null  int64
12   total_busy_partners                  194297 non-null  int64
13   total_outstanding_orders             194297 non-null  int64
14   created_at_hour                      194297 non-null  int64
15   created_at_date                     194297 non-null  object
16   delivery_time_min                    194297 non-null  float64
dtypes: datetime64[ns](2), float64(5), int64(9), object(1)
memory usage: 26.7+ MB
```

Extracting weekday from data

In [43]:

```
df['created_at_weekday']=df['created_at'].dt.weekday
```


In [44]:

```
df
```

Out[44]:

	market_id	created_at	actual_delivery_time	store_id	store_primary_category	order
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	50.395905		47.878758
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	48.408865		44.533674
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	48.408865		49.973629
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	48.408865		49.973629
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	48.408865		49.973629
...
197413	1.0	2015-02-17 00:19:41	2015-02-17 01:24:48	54.383918		44.039873
197414	1.0	2015-02-13 00:01:59	2015-02-13 00:58:22	54.383918		44.039873
197415	1.0	2015-01-24 04:46:08	2015-01-24 05:36:16	54.383918		44.039873
197416	1.0	2015-02-01 18:18:15	2015-02-01 19:23:22	48.883704		44.656187
197417	1.0	2015-02-08 19:24:33	2015-02-08 20:01:41	48.883704		44.656187

194297 rows × 18 columns




In [45]:

```
sns.kdeplot(df['delivery_time_min'])
```

Out[45]:

<AxesSubplot:xlabel='delivery_time_min', ylabel='Density'>




In [46]:

```
sns.boxplot(df['delivery_time_min'])
```


Out[46]:

<AxesSubplot:xlabel='delivery_time_min'>




In [47]:

```
sns.countplot(data=df, x='created_at_weekday', order=df['created_at_weekday'].value_count  
plt.xticks(rotation = 90)  
plt.show()  
# There more orders on weekday 5 and 6 i.e week day
```




In [48]:

```
sns.countplot(data=df, x='created_at_weekday', order=df['created_at_weekday'].value_count)
plt.xticks(rotation = 90)
plt.show()
# Least number of orders are on weekday 6
```




In [49]:

```
sns.countplot(data=df, x='created_at_hour', order=df['created_at_hour'].value_counts().in
plt.xticks(rotation = 90)
plt.show()
# Most of the orders are created at hours 2, 1 and 3
```



In [50]:

```
sns.countplot(data=df, x='created_at_hour', order=df['created_at_hour'].value_counts().in
plt.xticks(rotation = 90)
plt.show()
```



In [51]:

```
df.info()
```

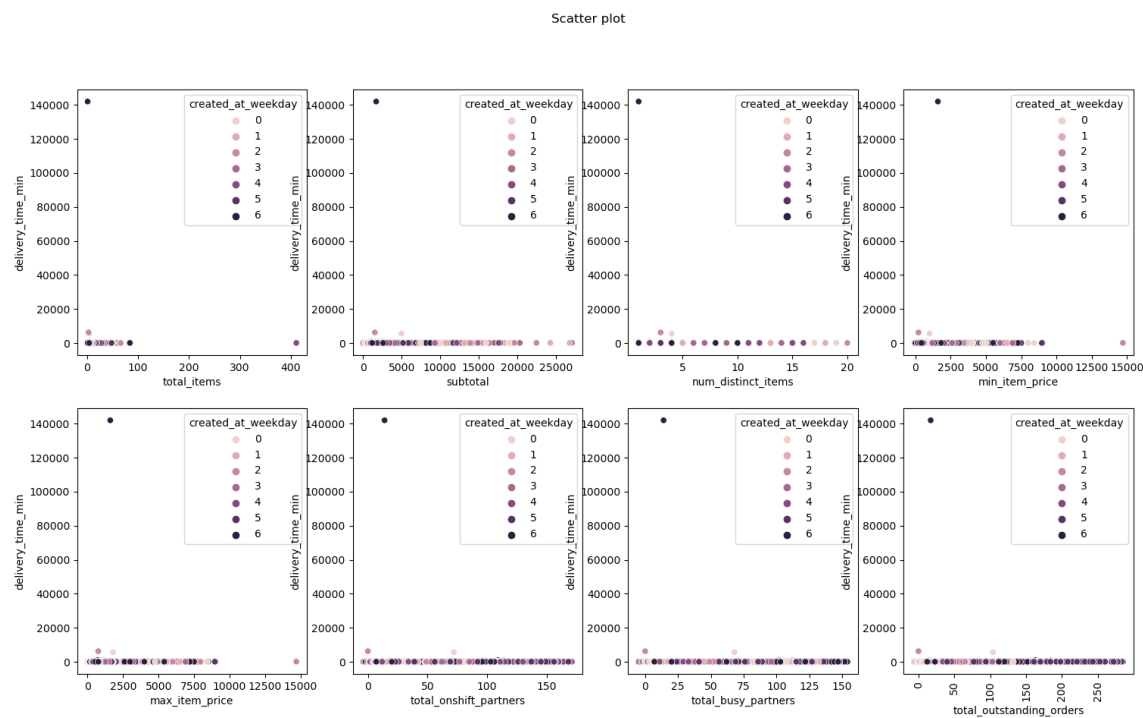
```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 194297 entries, 0 to 197417
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                             194297 non-null  float64
1   created_at                             194297 non-null  datetime64[ns]
2   actual_delivery_time                   194297 non-null  datetime64[ns]
3   store_id                               194297 non-null  float64
4   store_primary_category                 194297 non-null  float64
5   order_protocol                         194297 non-null  float64
6   total_items                            194297 non-null  int64
7   subtotal                               194297 non-null  int64
8   num_distinct_items                     194297 non-null  int64
9   min_item_price                         194297 non-null  int64
10  max_item_price                         194297 non-null  int64
11  total_onshift_partners                  194297 non-null  int64
12  total_busy_partners                     194297 non-null  int64
13  total_outstanding_orders                194297 non-null  int64
14  created_at_hour                         194297 non-null  int64
15  created_at_date                         194297 non-null  object
16  delivery_time_min                       194297 non-null  float64
17  created_at_weekday                      194297 non-null  int64
dtypes: datetime64[ns](2), float64(5), int64(10), object(1)
memory usage: 28.2+ MB
```

In [52]:

```
fig, axes = plt.subplots(2, 4, figsize=(18, 10))
plt.xticks(rotation = 90)
fig.suptitle('Scatter plot')
sns.scatterplot(ax=axes[0, 0], data=df, x='total_items', y='delivery_time_min', hue='created_at_weekday')
sns.scatterplot(ax=axes[0, 1], data=df, x='subtotal', y='delivery_time_min', hue='created_at_weekday')
sns.scatterplot(ax=axes[0, 2], data=df, x='num_distinct_items', y='delivery_time_min', hue='created_at_weekday')
sns.scatterplot(ax=axes[0, 3], data=df, x='min_item_price', y='delivery_time_min', hue='created_at_weekday')
sns.scatterplot(ax=axes[1, 0], data=df, x='max_item_price', y='delivery_time_min', hue='created_at_weekday')
sns.scatterplot(ax=axes[1, 1], data=df, x='total_onshift_partners', y='delivery_time_min', hue='created_at_weekday')
sns.scatterplot(ax=axes[1, 2], data=df, x='total_busy_partners', y='delivery_time_min', hue='created_at_weekday')
sns.scatterplot(ax=axes[1, 3], data=df, x='total_outstanding_orders', y='delivery_time_min', hue='created_at_weekday')
```

Out[52]:

<AxesSubplot:xlabel='total_outstanding_orders', ylabel='delivery_time_min'>



In [53]:

```
df['created_at_date_month'] = pd.to_datetime(df['created_at_date']).dt.month
df['created_at_date_year'] = pd.to_datetime(df['created_at_date']).dt.year
df['created_at_date_day'] = pd.to_datetime(df['created_at_date']).dt.day
# Extracting month year and day
```

In [54]:

df.columns

Out[54]:

```
Index(['market_id', 'created_at', 'actual_delivery_time', 'store_id',
      'store_primary_category', 'order_protocol', 'total_items', 'subtotal',
      'num_distinct_items', 'min_item_price', 'max_item_price',
      'total_onshift_partners', 'total_busy_partners',
      'total_outstanding_orders', 'created_at_hour', 'created_at_date',
      'delivery_time_min', 'created_at_weekday', 'created_at_date_month',
      'created_at_date_year', 'created_at_date_day'],
      dtype='object')
```

In [55]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 194297 entries, 0 to 197417
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                             194297 non-null float64
1   created_at                             194297 non-null datetime64[ns]
2   actual_delivery_time                   194297 non-null datetime64[ns]
3   store_id                               194297 non-null float64
4   store_primary_category                 194297 non-null float64
5   order_protocol                         194297 non-null float64
6   total_items                            194297 non-null int64
7   subtotal                               194297 non-null int64
8   num_distinct_items                     194297 non-null int64
9   min_item_price                         194297 non-null int64
10  max_item_price                         194297 non-null int64
11  total_onshift_partners                  194297 non-null int64
12  total_busy_partners                     194297 non-null int64
13  total_outstanding_orders                194297 non-null int64
14  created_at_hour                         194297 non-null int64
15  created_at_date                         194297 non-null object
16  delivery_time_min                       194297 non-null float64
17  created_at_weekday                      194297 non-null int64
18  created_at_date_month                   194297 non-null int64
19  created_at_date_year                    194297 non-null int64
20  created_at_date_day                     194297 non-null int64
dtypes: datetime64[ns](2), float64(5), int64(13), object(1)
memory usage: 32.6+ MB
```

In [56]:

```
dff=df[['market_id','store_id',
        'store_primary_category', 'order_protocol', 'total_items', 'subtotal',
        'num_distinct_items', 'min_item_price', 'max_item_price',
        'total_onshift_partners', 'total_busy_partners',
        'total_outstanding_orders', 'created_at_hour',
        'delivery_time_min', 'created_at_weekday', 'created_at_date_month',
        'created_at_date_year', 'created_at_date_day']]
```

In [57]:

```
dff.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 194297 entries, 0 to 197417
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                             194297 non-null float64
1   store_id                              194297 non-null float64
2   store_primary_category                 194297 non-null float64
3   order_protocol                         194297 non-null float64
4   total_items                           194297 non-null int64
5   subtotal                              194297 non-null int64
6   num_distinct_items                    194297 non-null int64
7   min_item_price                        194297 non-null int64
8   max_item_price                        194297 non-null int64
9   total_onshift_partners                 194297 non-null int64
10  total_busy_partners                    194297 non-null int64
11  total_outstanding_orders                194297 non-null int64
12  created_at_hour                        194297 non-null int64
13  delivery_time_min                      194297 non-null float64
14  created_at_weekday                     194297 non-null int64
15  created_at_date_month                   194297 non-null int64
16  created_at_date_year                   194297 non-null int64
17  created_at_date_day                     194297 non-null int64
dtypes: float64(5), int64(13)
memory usage: 28.2 MB
```

In [58]:

```
X=np.array(dff)
```

Outlier removal

In [59]:

```
clf = LocalOutlierFactor(n_neighbors=25,contamination=0.05)
```

In [60]:

```
Y=clf.fit_predict(X)
```

In [61]:

```
pd.Series(Y).value_counts()
```

Out[61]:

```
1    184582  
-1     9715  
dtype: int64
```

In [62]:

```
df['out']=pd.Series(Y)
```

In [63]:

```
df1=df[df['out']==1]  
# Removing 9715 rows
```

In [64]:

```
df1=df1[df1['delivery_time_min']<20000]
```

In [65]:

```

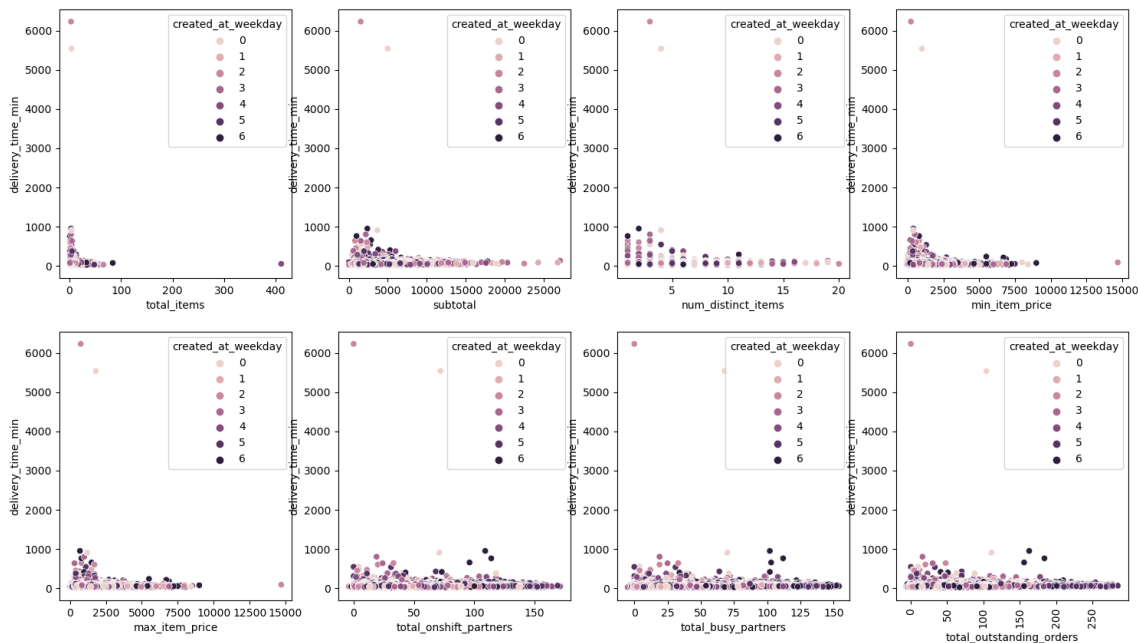
fig, axes = plt.subplots(2, 4, figsize=(18, 10))
plt.xticks(rotation = 90)
fig.suptitle('Scatter plot')
sns.scatterplot(ax=axes[0, 0], data=df1, x='total_items', y='delivery_time_min', hue='creat
sns.scatterplot(ax=axes[0, 1], data=df1, x='subtotal', y='delivery_time_min', hue='created_
sns.scatterplot(ax=axes[0, 2], data=df1, x='num_distinct_items', y='delivery_time_min', hue
sns.scatterplot(ax=axes[0, 3], data=df1, x='min_item_price', y='delivery_time_min', hue='cr
sns.scatterplot(ax=axes[1, 0], data=df1, x='max_item_price', y='delivery_time_min', hue='cr
sns.scatterplot(ax=axes[1, 1], data=df1, x='total_onshift_partners', y='delivery_time_min'
sns.scatterplot(ax=axes[1, 2], data=df1, x='total_busy_partners', y='delivery_time_min', hu
sns.scatterplot(ax=axes[1, 3], data=df1, x='total_outstanding_orders', y='delivery_time_mi

```

Out[65]:

<AxesSubplot:xlabel='total_outstanding_orders', ylabel='delivery_time_min'>

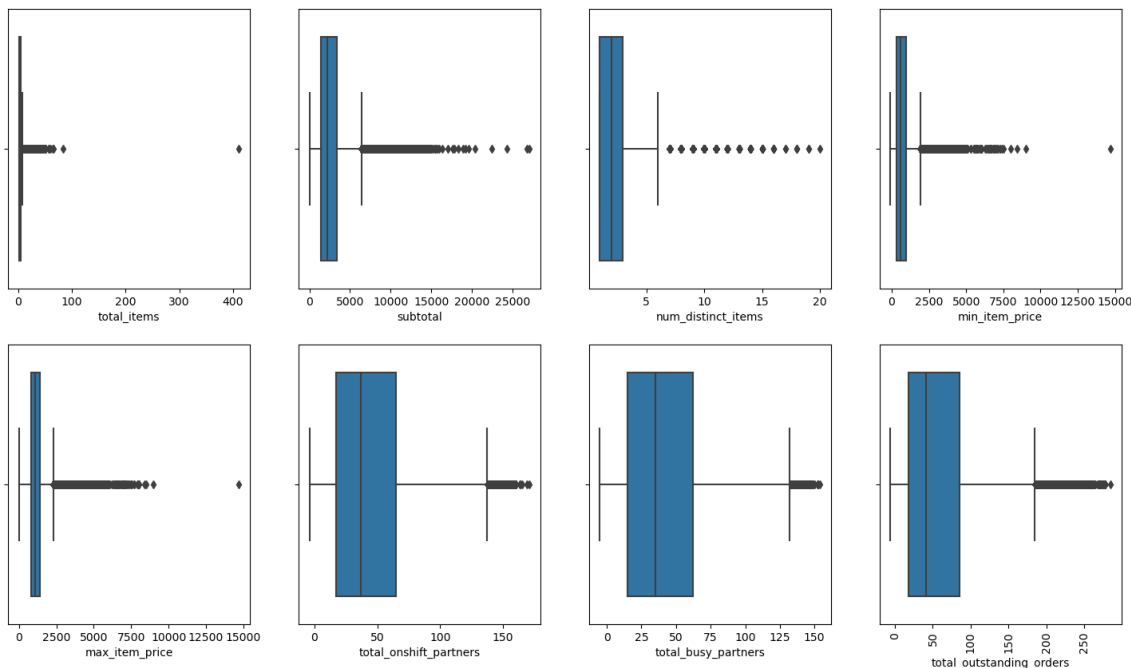
Scatter plot



In [66]:

```
fig, axes = plt.subplots(2, 4, figsize=(18, 10))
plt.xticks(rotation = 90)
fig.suptitle('Box plot')
sns.boxplot(ax=axes[0, 0], data=df1, x='total_items')
sns.boxplot(ax=axes[0, 1], data=df1, x='subtotal')
sns.boxplot(ax=axes[0, 2], data=df1, x='num_distinct_items' )
sns.boxplot(ax=axes[0, 3], data=df1, x='min_item_price' )
sns.boxplot(ax=axes[1, 0], data=df1, x='max_item_price')
sns.boxplot(ax=axes[1, 1], data=df1, x='total_onshift_partners')
sns.boxplot(ax=axes[1, 2], data=df1, x='total_busy_partners')
sns.boxplot(ax=axes[1, 3], data=df1, x='total_outstanding_orders')
plt.show()
```

Box plot



In [67]:

```
df1.columns
```

Out[67]:

```
Index(['market_id', 'created_at', 'actual_delivery_time', 'store_id',
      'store_primary_category', 'order_protocol', 'total_items', 'subtotal',
      'num_distinct_items', 'min_item_price', 'max_item_price',
      'total_onshift_partners', 'total_busy_partners',
      'total_outstanding_orders', 'created_at_hour', 'created_at_date',
      'delivery_time_min', 'created_at_weekday', 'created_at_date_month',
      'created_at_date_year', 'created_at_date_day', 'out'],
      dtype='object')
```

In [68]:

df1.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 181663 entries, 0 to 194296
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            181663 non-null  float64
1   created_at                           181663 non-null  datetime64[ns]
2   actual_delivery_time                 181663 non-null  datetime64[ns]
3   store_id                             181663 non-null  float64
4   store_primary_category               181663 non-null  float64
5   order_protocol                       181663 non-null  float64
6   total_items                          181663 non-null  int64
7   subtotal                             181663 non-null  int64
8   num_distinct_items                   181663 non-null  int64
9   min_item_price                       181663 non-null  int64
10  max_item_price                       181663 non-null  int64
11  total_onshift_partners               181663 non-null  int64
12  total_busy_partners                  181663 non-null  int64
13  total_outstanding_orders             181663 non-null  int64
14  created_at_hour                      181663 non-null  int64
15  created_at_date                      181663 non-null  object
16  delivery_time_min                    181663 non-null  float64
17  created_at_weekday                   181663 non-null  int64
18  created_at_date_month                181663 non-null  int64
19  created_at_date_year                 181663 non-null  int64
20  created_at_date_day                  181663 non-null  int64
21  out                                  181663 non-null  float64
dtypes: datetime64[ns](2), float64(6), int64(13), object(1)
memory usage: 31.9+ MB
```

In [69]:

```
X=df1[['market_id', 'store_id',
       'store_primary_category', 'order_protocol', 'total_items', 'subtotal',
       'num_distinct_items', 'min_item_price', 'max_item_price',
       'total_onshift_partners', 'total_busy_partners',
       'total_outstanding_orders', 'created_at_hour', 'created_at_weekday', 'created_at',
       'created_at_date_year', 'created_at_date_day']]
y=df1['delivery_time_min']
```


In [70]:

```
X
```

Out[70]:

	market_id	store_id	store_primary_category	order_protocol	total_items	subtotal
0	1.0	50.395905	47.878758	1.0	4	3441
1	2.0	48.408865	44.533674	2.0	1	1900
2	3.0	48.408865	49.973629	1.0	1	1900
3	3.0	48.408865	49.973629	1.0	6	6900
4	3.0	48.408865	49.973629	1.0	3	3900
...
194292	6.0	46.855428	44.921455	2.0	1	1725
194293	6.0	46.855428	44.921455	2.0	1	1225
194294	6.0	46.855428	44.921455	2.0	1	1325
194295	1.0	46.855428	47.402817	3.0	1	1925
194296	6.0	46.855428	44.921455	2.0	1	1925

181663 rows × 17 columns

Train test split

In [71]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
X_train_val, X_train, y_train_val, y_train = train_test_split(X_train, y_train, test_siz
```

In [72]:

```
len(X_train), len(X_train_val), len(X_test)
```

Out[72]:

(147146, 16350, 18167)

Random Forest

In [74]:

```
model=RandomForestRegressor(n_estimators=100, criterion='squared_error',
                             min_samples_split=2, min_samples_leaf=1, random_state=42)
```

In [75]:

```
model.fit(X_train,y_train)
```

Out[75]:

```
RandomForestRegressor(random_state=42)
```

In [76]:

```
model.score(X_train, y_train)
```

Out[76]:

```
0.8839544982478807
```

In [77]:

```
model.score(X_train_val, y_train_val)
```

Out[77]:

```
0.3060957896075144
```

In [78]:

```
model.score(X_test,y_test)
```

Out[78]:

```
0.04430276036233449
```

In [79]:

```
y_train_pred=model.predict(X_train)
y_train_val_pred=model.predict(X_train_val)
y_test_pred=model.predict(X_test)
```

In [80]:

```
mae = mean_absolute_error(y_train, y_train_pred)
mse = mean_squared_error(y_train, y_train_pred)
rmse = np.sqrt(mse)
```

In [81]:

```
print(mae,mse,rmse)
```

```
4.077480348256855 66.88168821794162 8.178122536251314
```

In [82]:

```
mae = mean_absolute_error(y_train_val, y_train_val_pred)
mse = mean_squared_error(y_train_val, y_train_val_pred)
rmse = np.sqrt(mse)
```

In [83]:

```
print(mae,mse,rmse)
```

```
11.01597553715839 263.8691448959443 16.244049522700436
```

In [84]:

```
mae = mean_absolute_error(y_test, y_test_pred)
mse = mean_squared_error(y_test, y_test_pred)
rmse = np.sqrt(mse)
```

In [85]:

```
print(mae,mse,rmse)
```

```
11.343540791213092 2429.767955019698 49.292676484643216
```

In [86]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
```

Data Scaling

In [87]:

```
scaler=StandardScaler()
X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
```

In [88]:

```
X_train, X_train_val, y_train, y_train_val = train_test_split(X_train, y_train, test_size=0.1, random_state=42)
```

In [89]:

```
X_train.shape
```

Out[89]:

```
(147146, 17)
```

In [90]:

```
X_train_val.shape
```

Out[90]:

```
(16350, 17)
```

In [91]:

```
X_test.shape
```

Out[91]:

```
(18167, 17)
```

Neural Network

In [1]:

```
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.layers import Activation
from tensorflow.keras.activations import linear
from tensorflow.keras import Sequential
import tensorflow as tf
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Dropout
from tensorflow.keras.callbacks import LearningRateScheduler
```

In [2]:

```
def scheduler(epoch, lr):
    if epoch < 400:
        return lr
    else:
        return lr * tf.math.exp(-0.1)
```

In [3]:

```
L2Reg = tf.keras.regularizers.L1L2(l1=1e-7, l2=1e-7)
callback = tf.keras.callbacks.LearningRateScheduler(scheduler)
```


In [4]:

```
model = Sequential([
    Dense(32, activation="relu", input_shape=(17,), name="hidden_1",kernel_regularizer = L2
    Dense(64, activation="relu", name="hidden_2",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(128, activation="relu", name="hidden_3",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(256, activation="relu", name="hidden_4",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(512, activation="relu", name="hidden_5",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(1024, activation="relu", name="hidden_6",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(2028, activation="relu", name="hidden_7",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(4056, activation="relu", name="hidden_8",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(8052, activation="relu", name="hidden_9",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(4056, activation="relu", name="hidden_10",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(2028, activation="relu", name="hidden_11",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(1024, activation="relu", name="hidden_12",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(512, activation="relu", name="hidden_13",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(256, activation="relu", name="hidden_14",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),

    Dense(128, activation="relu", name="hidden_15",kernel_regularizer = L2Reg ),
    BatchNormalization(),
    Dropout(0.3),
```

```
Dense(64, activation="relu", name="hidden_16", kernel_regularizer = L2Reg ),
BatchNormalization(),
Dropout(0.3),

Dense(32, activation="relu", name="hidden_17", kernel_regularizer = L2Reg ),
BatchNormalization(),
Dropout(0.3),

Dense(16, activation="relu", name="hidden_18", kernel_regularizer = L2Reg ),
BatchNormalization(),
Dropout(0.2),

Dense(8, activation="relu", name="hidden_19", kernel_regularizer = L2Reg ),
BatchNormalization(),
Dropout(0.2),

Dense(4, activation="relu", name="hidden_20", kernel_regularizer = L2Reg ),
BatchNormalization(),
Dropout(0.2),

Dense(2, activation="relu", name="hidden_21", kernel_regularizer = L2Reg ),
BatchNormalization(),

Dense(1, activation="linear", name="output")
])
```

In [5]:

```
model.summary()
```


Model: "sequential"

Layer (type)	Output Shape	Param #
hidden_1 (Dense)	(None, 32)	576
hidden_2 (Dense)	(None, 64)	2112
batch_normalization (Batch Normalization)	(None, 64)	256
dropout (Dropout)	(None, 64)	0
hidden_3 (Dense)	(None, 128)	8320
batch_normalization_1 (Batch Normalization)	(None, 128)	512
dropout_1 (Dropout)	(None, 128)	0
hidden_4 (Dense)	(None, 256)	33024
batch_normalization_2 (Batch Normalization)	(None, 256)	1024
dropout_2 (Dropout)	(None, 256)	0
hidden_5 (Dense)	(None, 512)	131584
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_3 (Dropout)	(None, 512)	0
hidden_6 (Dense)	(None, 1024)	525312
batch_normalization_4 (Batch Normalization)	(None, 1024)	4096
dropout_4 (Dropout)	(None, 1024)	0
hidden_7 (Dense)	(None, 2028)	2078700
batch_normalization_5 (Batch Normalization)	(None, 2028)	8112
dropout_5 (Dropout)	(None, 2028)	0
hidden_8 (Dense)	(None, 4056)	8229624
batch_normalization_6 (Batch Normalization)	(None, 4056)	16224
dropout_6 (Dropout)	(None, 4056)	0
hidden_9 (Dense)	(None, 8052)	32666964
batch_normalization_7 (Batch Normalization)	(None, 8052)	32208
dropout_7 (Dropout)	(None, 8052)	0

hidden_10 (Dense)	(None, 4056)	32662968
batch_normalization_8 (Batch Normalization)	(None, 4056)	16224
dropout_8 (Dropout)	(None, 4056)	0
hidden_11 (Dense)	(None, 2028)	8227596
batch_normalization_9 (Batch Normalization)	(None, 2028)	8112
dropout_9 (Dropout)	(None, 2028)	0
hidden_12 (Dense)	(None, 1024)	2077696
batch_normalization_10 (Batch Normalization)	(None, 1024)	4096
dropout_10 (Dropout)	(None, 1024)	0
hidden_13 (Dense)	(None, 512)	524800
batch_normalization_11 (Batch Normalization)	(None, 512)	2048
dropout_11 (Dropout)	(None, 512)	0
hidden_14 (Dense)	(None, 256)	131328
batch_normalization_12 (Batch Normalization)	(None, 256)	1024
dropout_12 (Dropout)	(None, 256)	0
hidden_15 (Dense)	(None, 128)	32896
batch_normalization_13 (Batch Normalization)	(None, 128)	512
dropout_13 (Dropout)	(None, 128)	0
hidden_16 (Dense)	(None, 64)	8256
batch_normalization_14 (Batch Normalization)	(None, 64)	256
dropout_14 (Dropout)	(None, 64)	0
hidden_17 (Dense)	(None, 32)	2080
batch_normalization_15 (Batch Normalization)	(None, 32)	128
dropout_15 (Dropout)	(None, 32)	0
hidden_18 (Dense)	(None, 16)	528
batch_normalization_16 (Batch Normalization)	(None, 16)	64

dropout_16 (Dropout)	(None, 16)	0
hidden_19 (Dense)	(None, 8)	136
batch_normalization_17 (Batch Normalization)	(None, 8)	32
dropout_17 (Dropout)	(None, 8)	0
hidden_20 (Dense)	(None, 4)	36
batch_normalization_18 (Batch Normalization)	(None, 4)	16
dropout_18 (Dropout)	(None, 4)	0
hidden_21 (Dense)	(None, 2)	10
batch_normalization_19 (Batch Normalization)	(None, 2)	8
output (Dense)	(None, 1)	3

```
=====
Total params: 87,441,549
Trainable params: 87,393,049
Non-trainable params: 48,500
```

In [6]:

```
from keras.utils.vis_utils import plot_model
plot_model(model, show_shapes=True, show_layer_names=True)
```

You must install pydot (`pip install pydot`) and install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) (<https://graphviz.gitlab.io/download/>) for plot_model to work.

In [118]:

```
model.compile(optimizer = tf.keras.optimizers.Adam(beta_1=0.9,beta_2=0.999),
              loss = tf.keras.losses.MeanAbsolutePercentageError(),
              metrics=['mean_absolute_error','mean_absolute_percentage_error','mean_squared_error'])
```

In [119]:

```
history = model.fit(X_train, y_train, validation_data = (X_train_val, y_train_val), epoch
```

Epoch 1/900

```
50/50 [=====] - 32s 274ms/step - loss: 99.9924
- mean_absolute_error: 47.7042 - mean_absolute_percentage_error: 99.872
2 - mean_squared_error: 2853.6221 - val_loss: 99.5867 - val_mean_absolu
te_error: 47.8246 - val_mean_absolute_percentage_error: 99.4633 - val_m
ean_squared_error: 2667.5596 - lr: 0.0010
```

Epoch 2/900

```
50/50 [=====] - 13s 258ms/step - loss: 99.6877
- mean_absolute_error: 47.5981 - mean_absolute_percentage_error: 99.563
7 - mean_squared_error: 2845.6909 - val_loss: 99.6468 - val_mean_absolu
te_error: 47.8593 - val_mean_absolute_percentage_error: 99.5224 - val_m
ean_squared_error: 2671.8210 - lr: 0.0010
```

Epoch 3/900


```
50/50 [=====] - 13s 259ms/step - loss: 99.2764
- mean_absolute_error: 47.4694 - mean_absolute_percentage_error: 99.151
9 - mean_squared_error: 2837.8337 - val_loss: 98.2855 - val_mean_absolu
te_error: 47.4054 - val_mean_absolute_percentage_error: 98.1610 - val_m
ean_squared_error: 2642.1807 - lr: 0.0010
```

Epoch 4/900

```
50/50 [=====] - 13s 258ms/step - loss: 98.7607
```


In [120]:

```
epochs = history.epoch
loss = history.history["loss"]
val_loss = history.history["val_loss"]
plt.plot(epochs, loss, label="train")
plt.plot(epochs, val_loss, label="validation")
plt.legend()
plt.title("Loss VS Epochs")
plt.show()
```



In [121]:

```
epochs = history.epoch[500:]
loss = history.history["loss"][500:]
val_loss = history.history["val_loss"][500:]
plt.plot(epochs, loss, label="train")
plt.plot(epochs, val_loss, label="validation")
plt.legend()
plt.title("Loss VS Epochs")
plt.show()
```



In [122]:

```
model.evaluate(X_train_val, y_train_val)
```

511/511 [=====] - 4s 8ms/step - loss: 22.0249 - mean_absolute_error: 11.1394 - mean_absolute_percentage_error: 21.9483 - mean_squared_error: 299.5685

Out[122]:

[22.024900436401367, 11.13938045501709, 21.948274612426758, 299.5685119628906]

In [123]:

```
model.evaluate(X_train, y_train)
```

4599/4599 [=====] - 36s 8ms/step - loss: 21.5963 - mean_absolute_error: 10.8487 - mean_absolute_percentage_error: 21.5198 - mean_squared_error: 489.0732

Out[123]:

[21.5963134765625, 10.848658561706543, 21.5197811126709, 489.0732116699219]

In [124]:

```
model.evaluate(X_test, y_test)
```

568/568 [=====] - 5s 8ms/step - loss: 22.0226 - mean_absolute_error: 11.4178 - mean_absolute_percentage_error: 21.9460 - mean_squared_error: 2468.3508

Out[124]:

[22.022571563720703, 11.417793273925781, 21.945960998535156, 2468.350830078125]

In []:

In []: